

# CONVOLUTIONAL NEURAL NETWORKS

COMP4660/8420

Bio-inspired Computing: Applications and Interfaces  
Adapted from lecture notes by Christopher Chow and  
Josephine Plested

# OVERVIEW

Biological inspirations

Structures and layers


Implementation and  
examples



# CONVOLUTIONAL NEURAL NETWORKS

- CNN or ConvNet
- Successful in spatial problem domains in particular, but increasingly in temporal domains too
- Prominent component of deep learning history
- Common and popular form of deep learning

# INSPIRED BY THE VISUAL PATHWAY



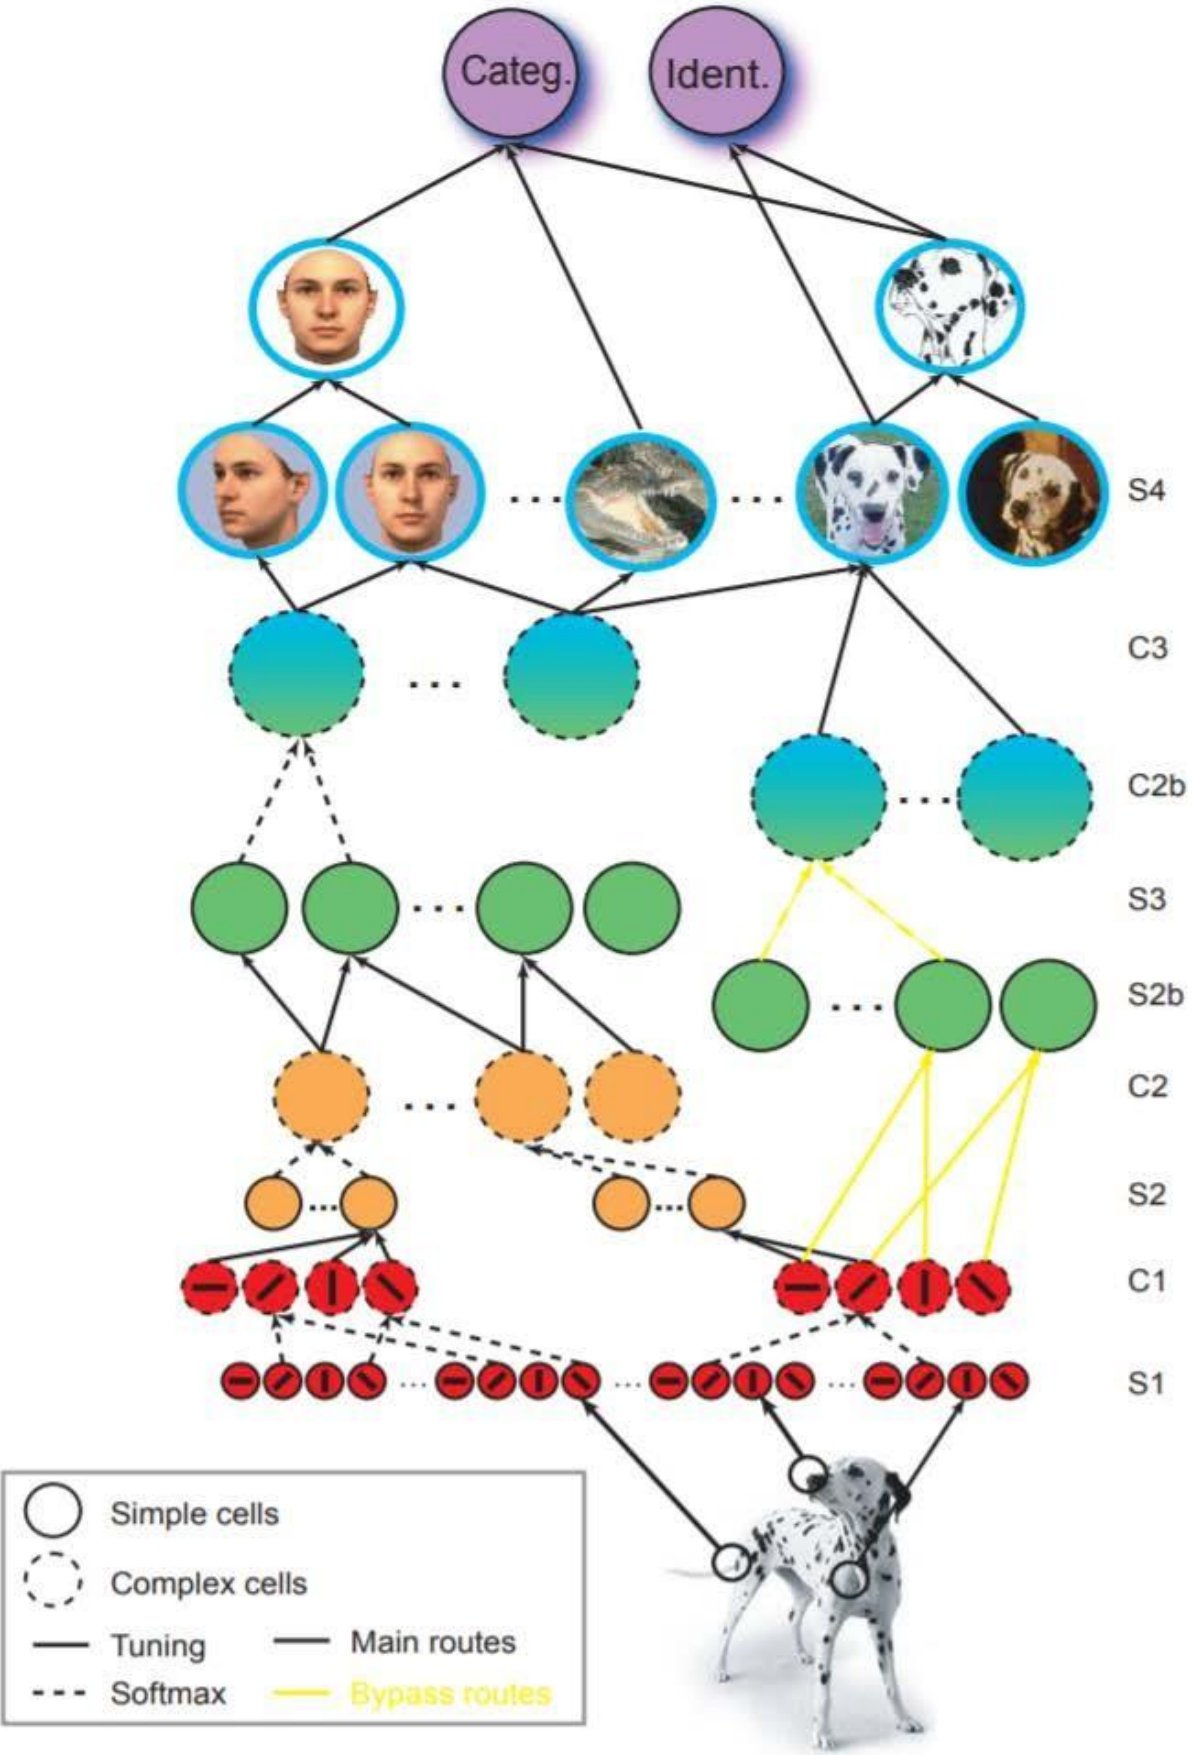
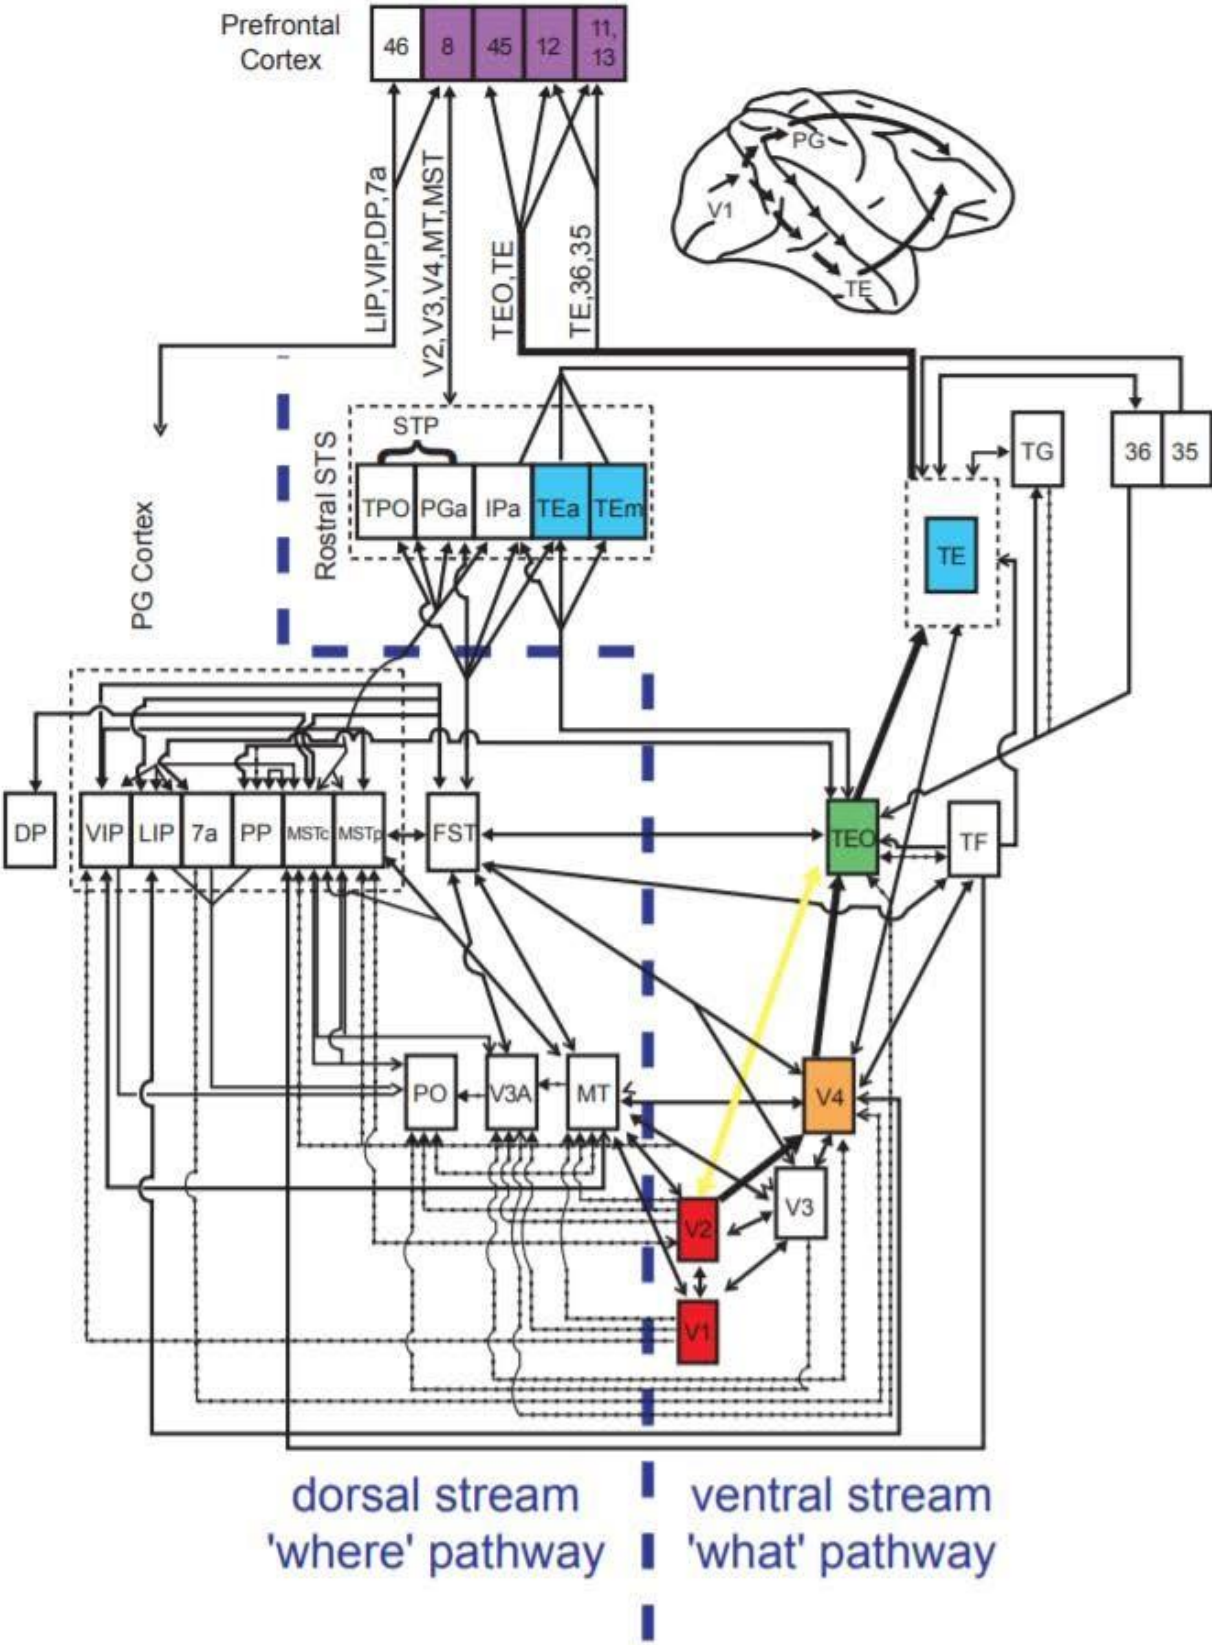
# VISUAL CORTEX VENTRAL PATHWAY

## LGN-V1-V2-V4-IT hierarchy

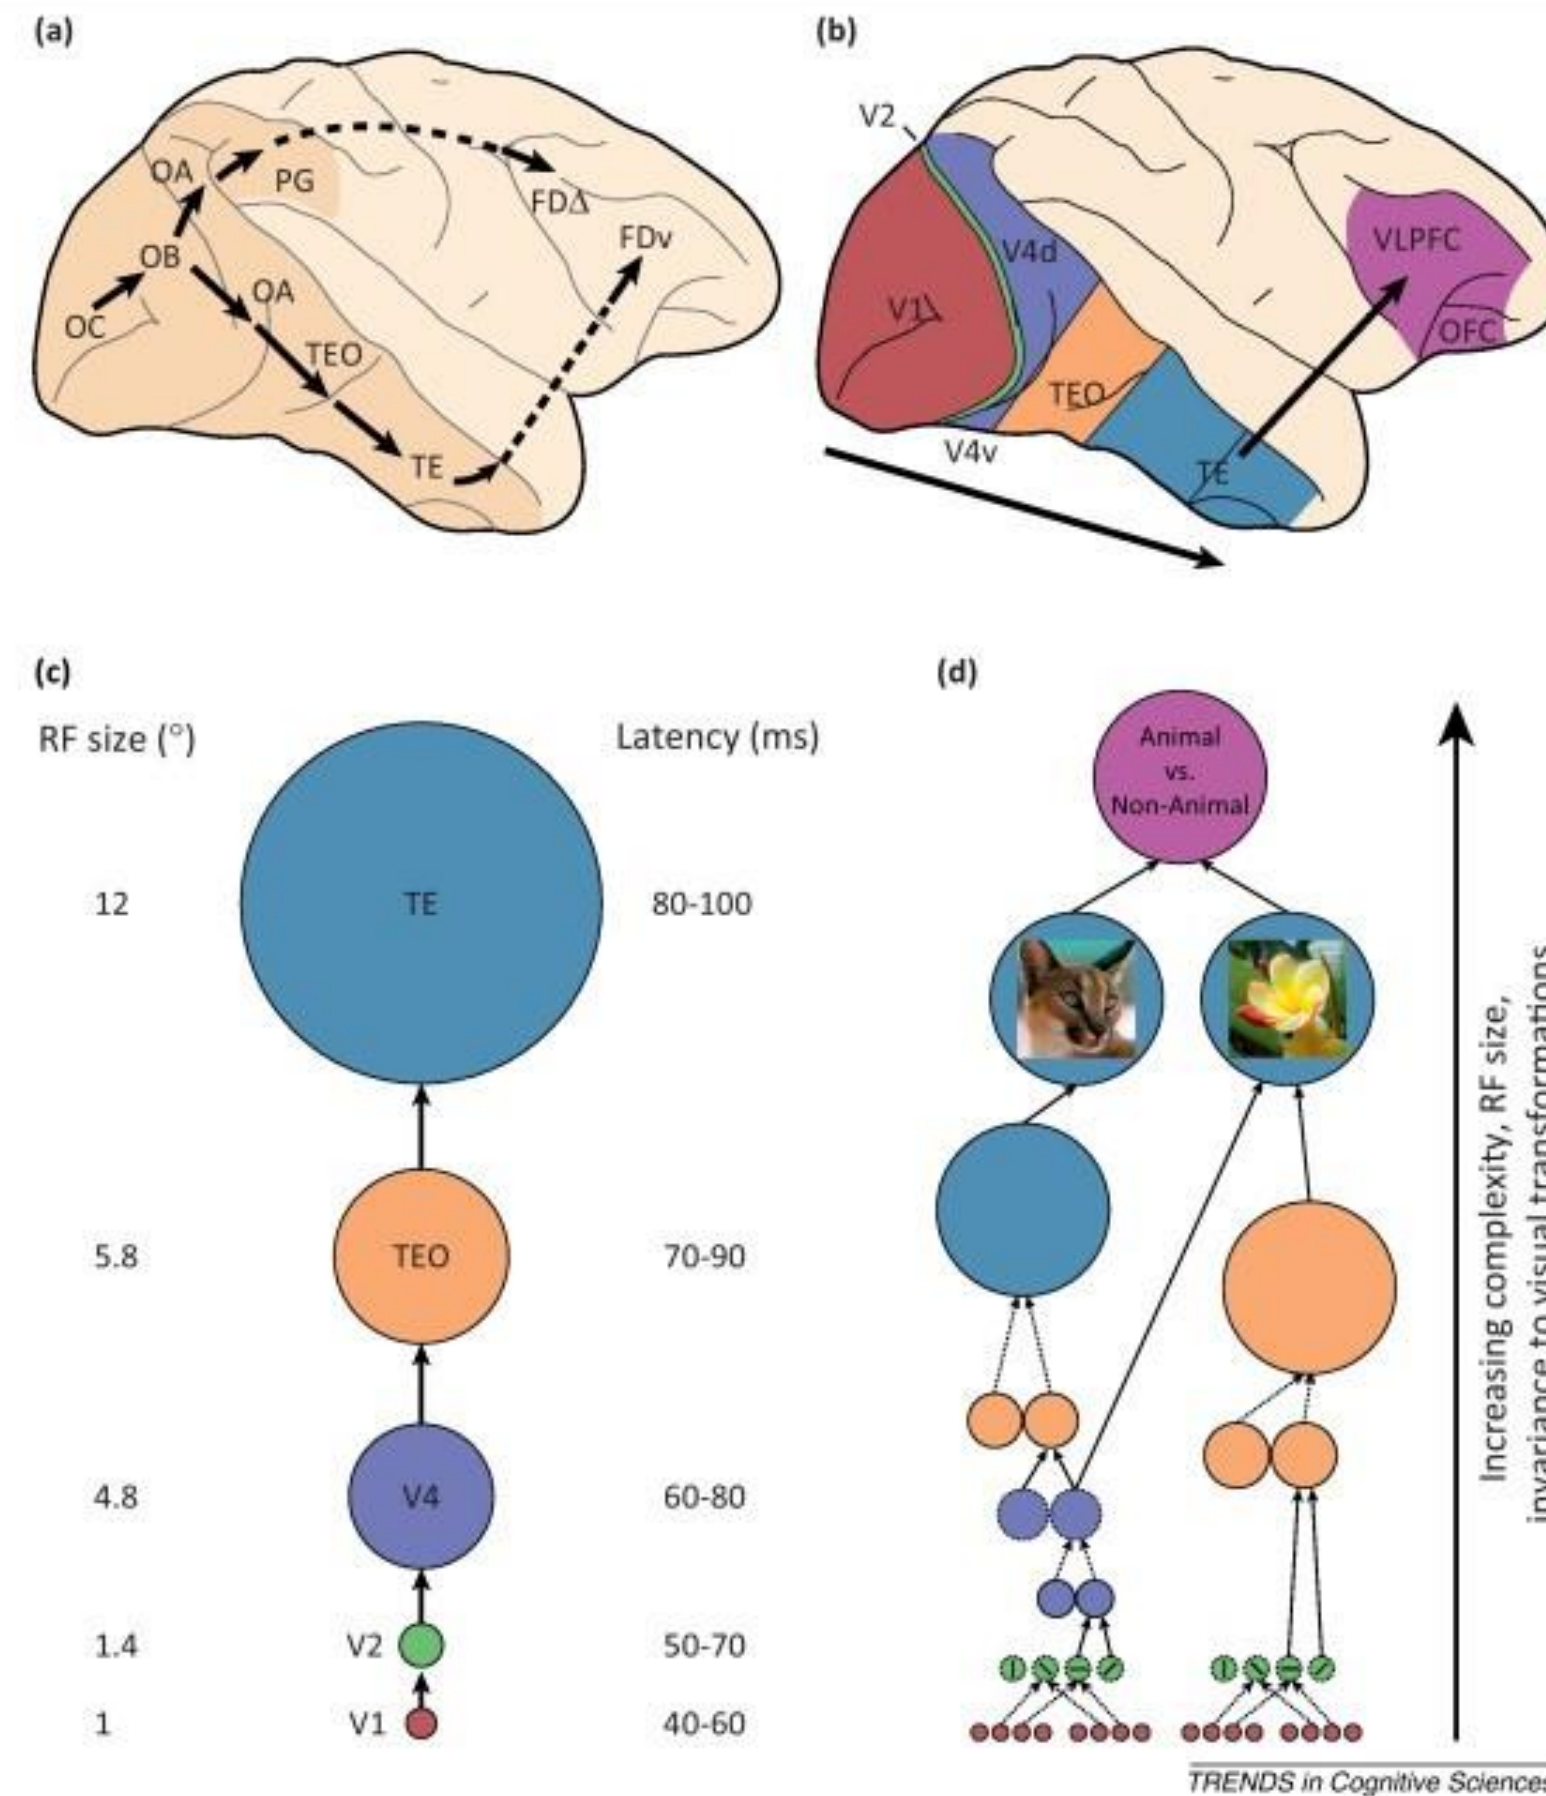
- Lateral geniculate nucleus (LGN)
  - Provides sensory input from retina to brain (visual cortex)
  - Calculates spatial dimensions
- Visual area one of visual cortex (V1)
  - Neurons with similar tuning properties cluster together
  - Edge detection
- Visual area two of visual cortex (V2)
  - Complex properties such as contours and orientation
- Visual area four of visual cortex (V4)
  - More complex properties such as shapes, spatial frequency, colour
- Inferior temporal gyrus (IT)
  - Tuned for very complex properties such as faces, objects, patterns



# HIERARCHICAL VISUAL PROCESSING



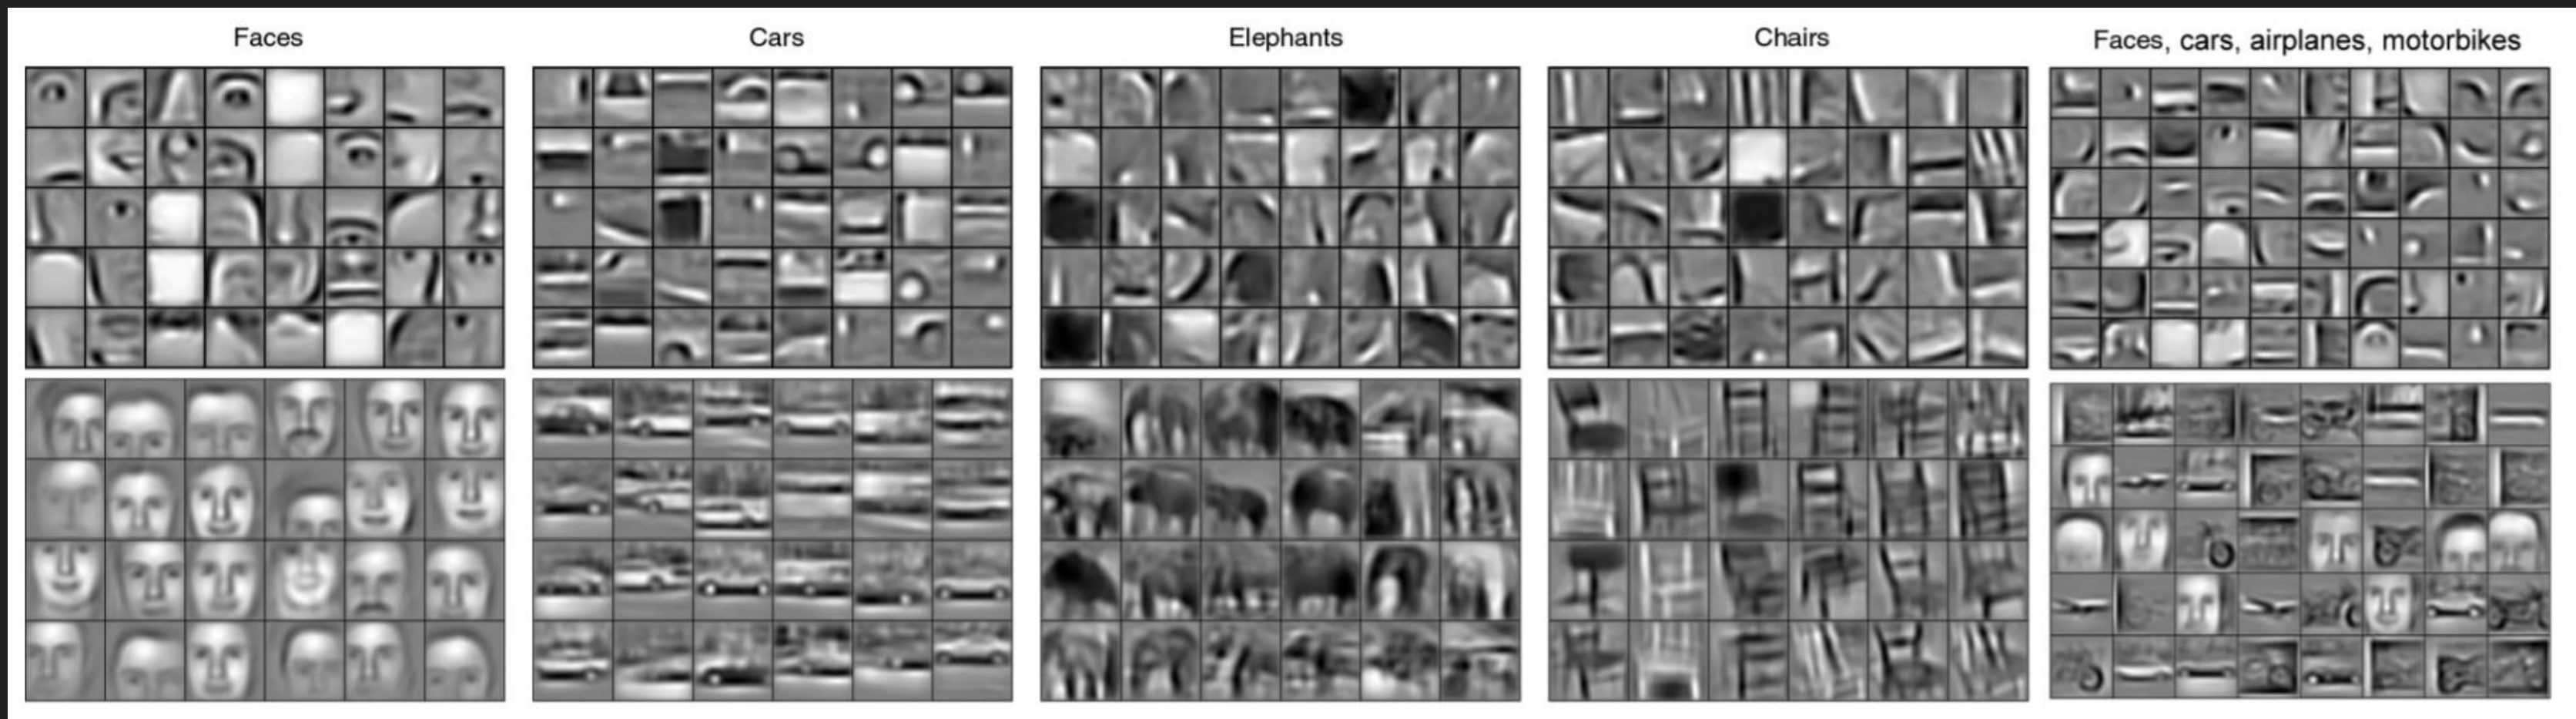
# INCREASING RECEPTIVE FIELD SIZE






## GENERAL INTUITION

- Transforms **raw input**, layer-by-layer, to final class scores
- More complex features are learnt as you go deeper into the network by combining simpler ones
- Higher level abstracted representations are combinations of several lower level features



# INTUITION AND INSPIRATION

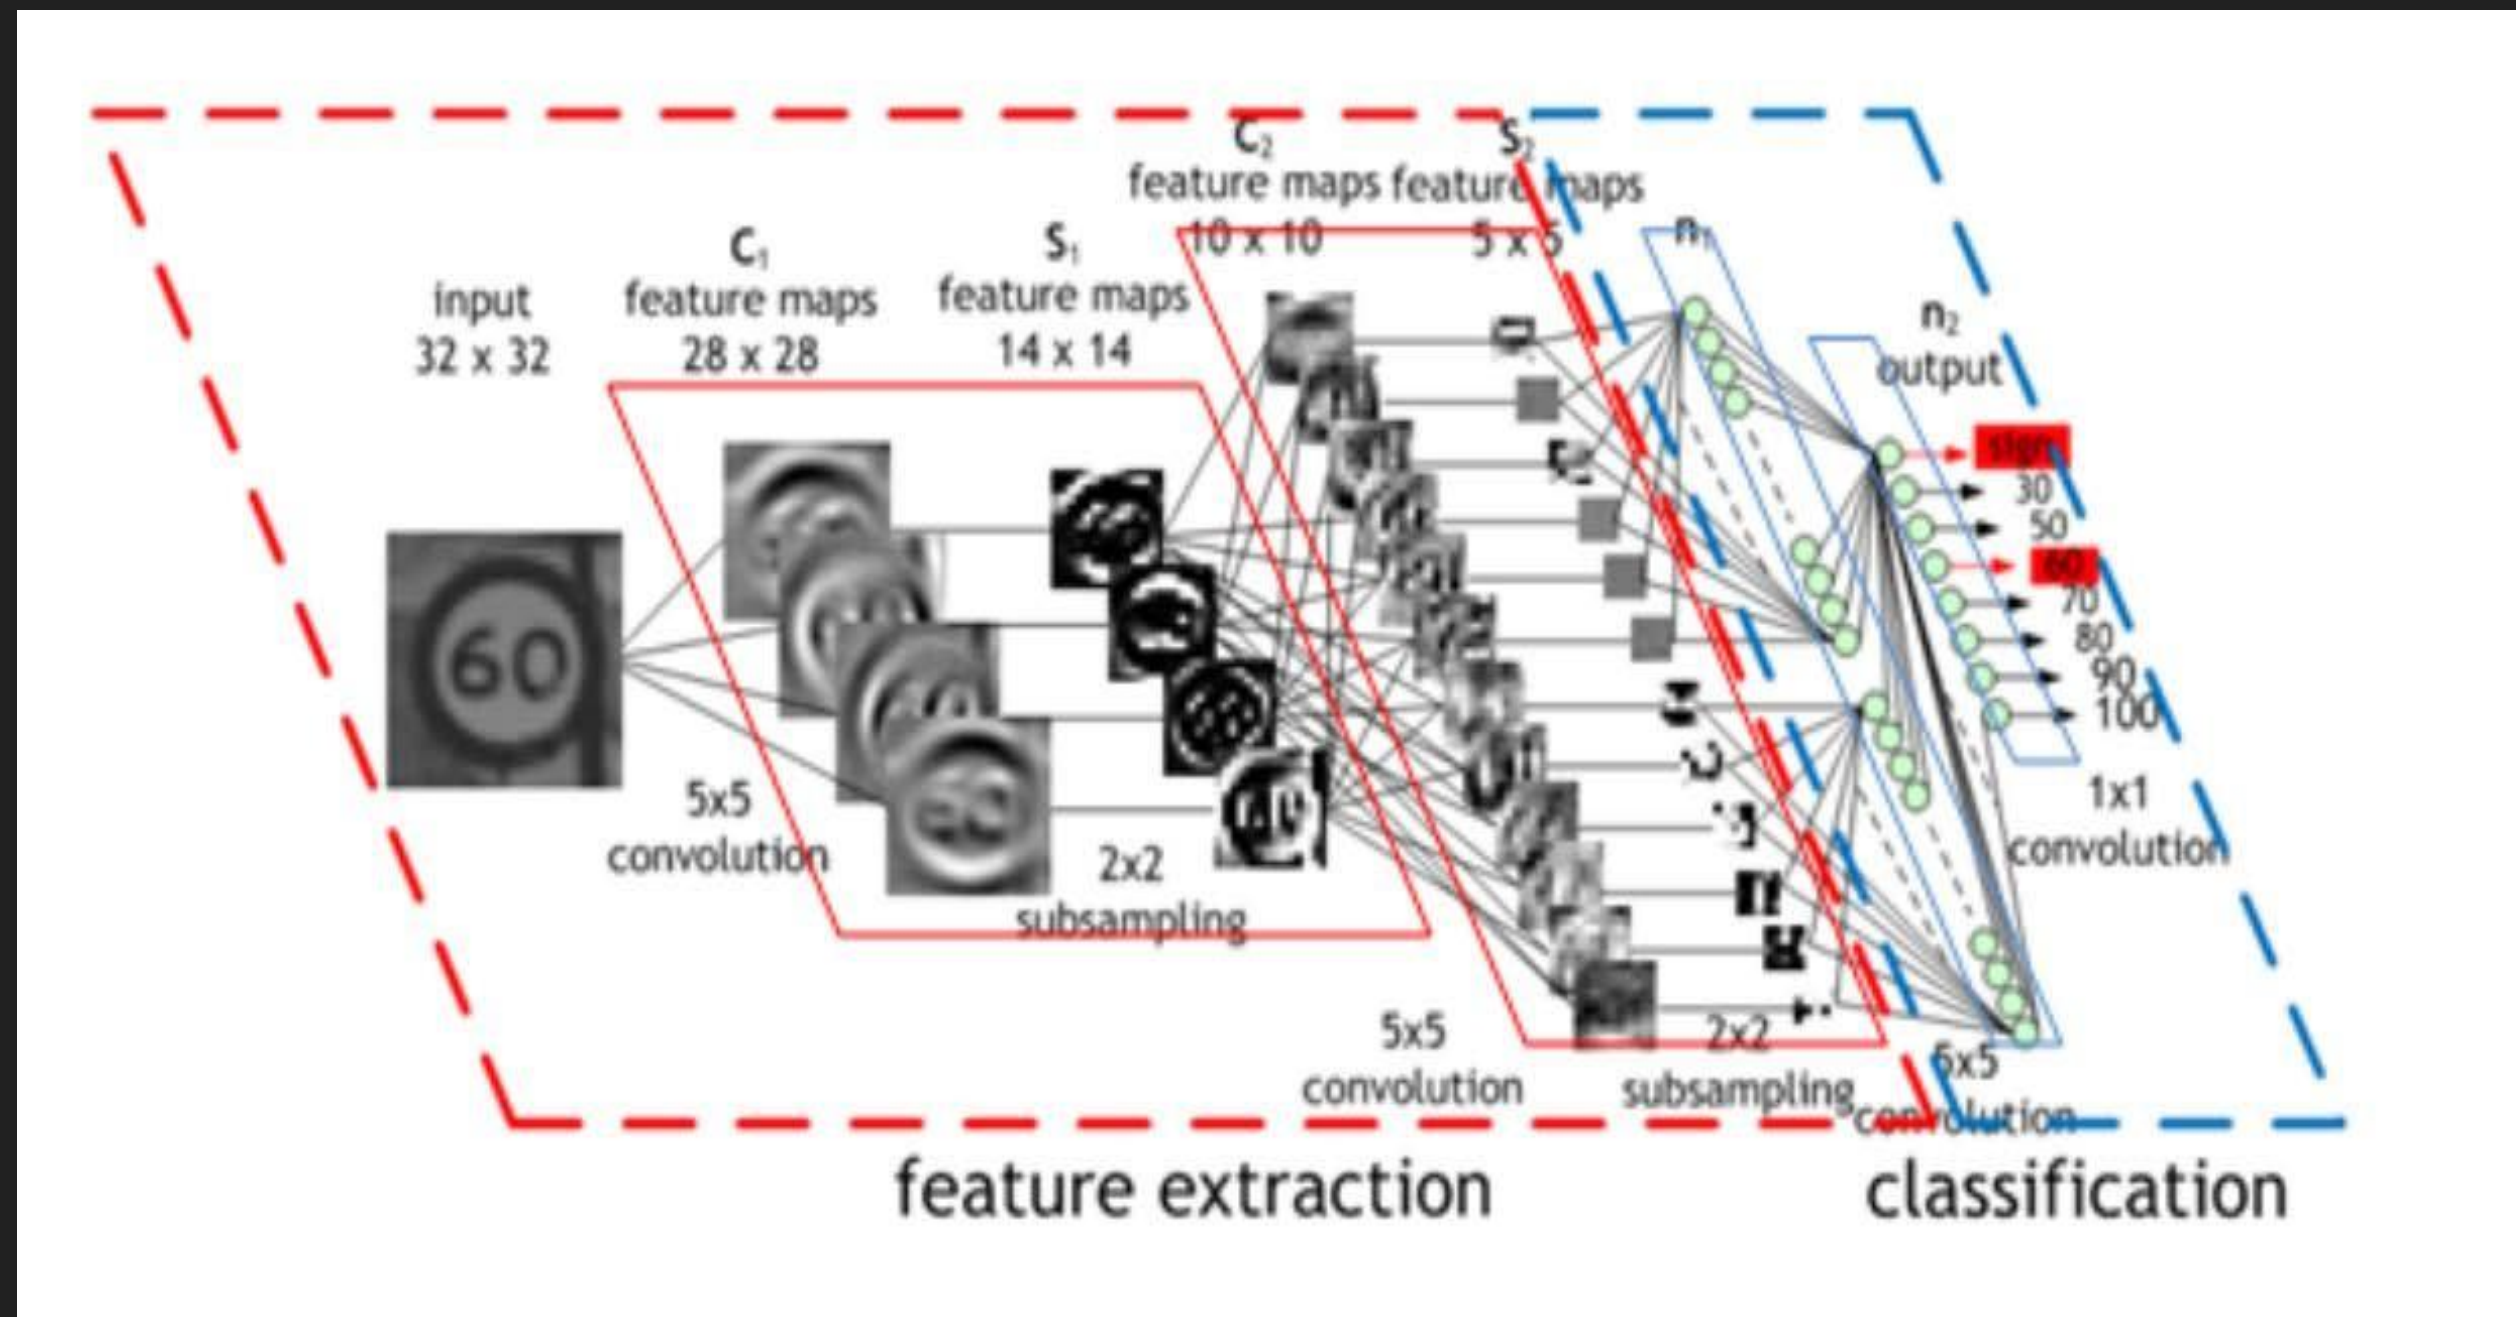
Specialises in processing **grid-like** data



Uses **sparse interactions** and **parameter sharing** to achieve **equivariance**

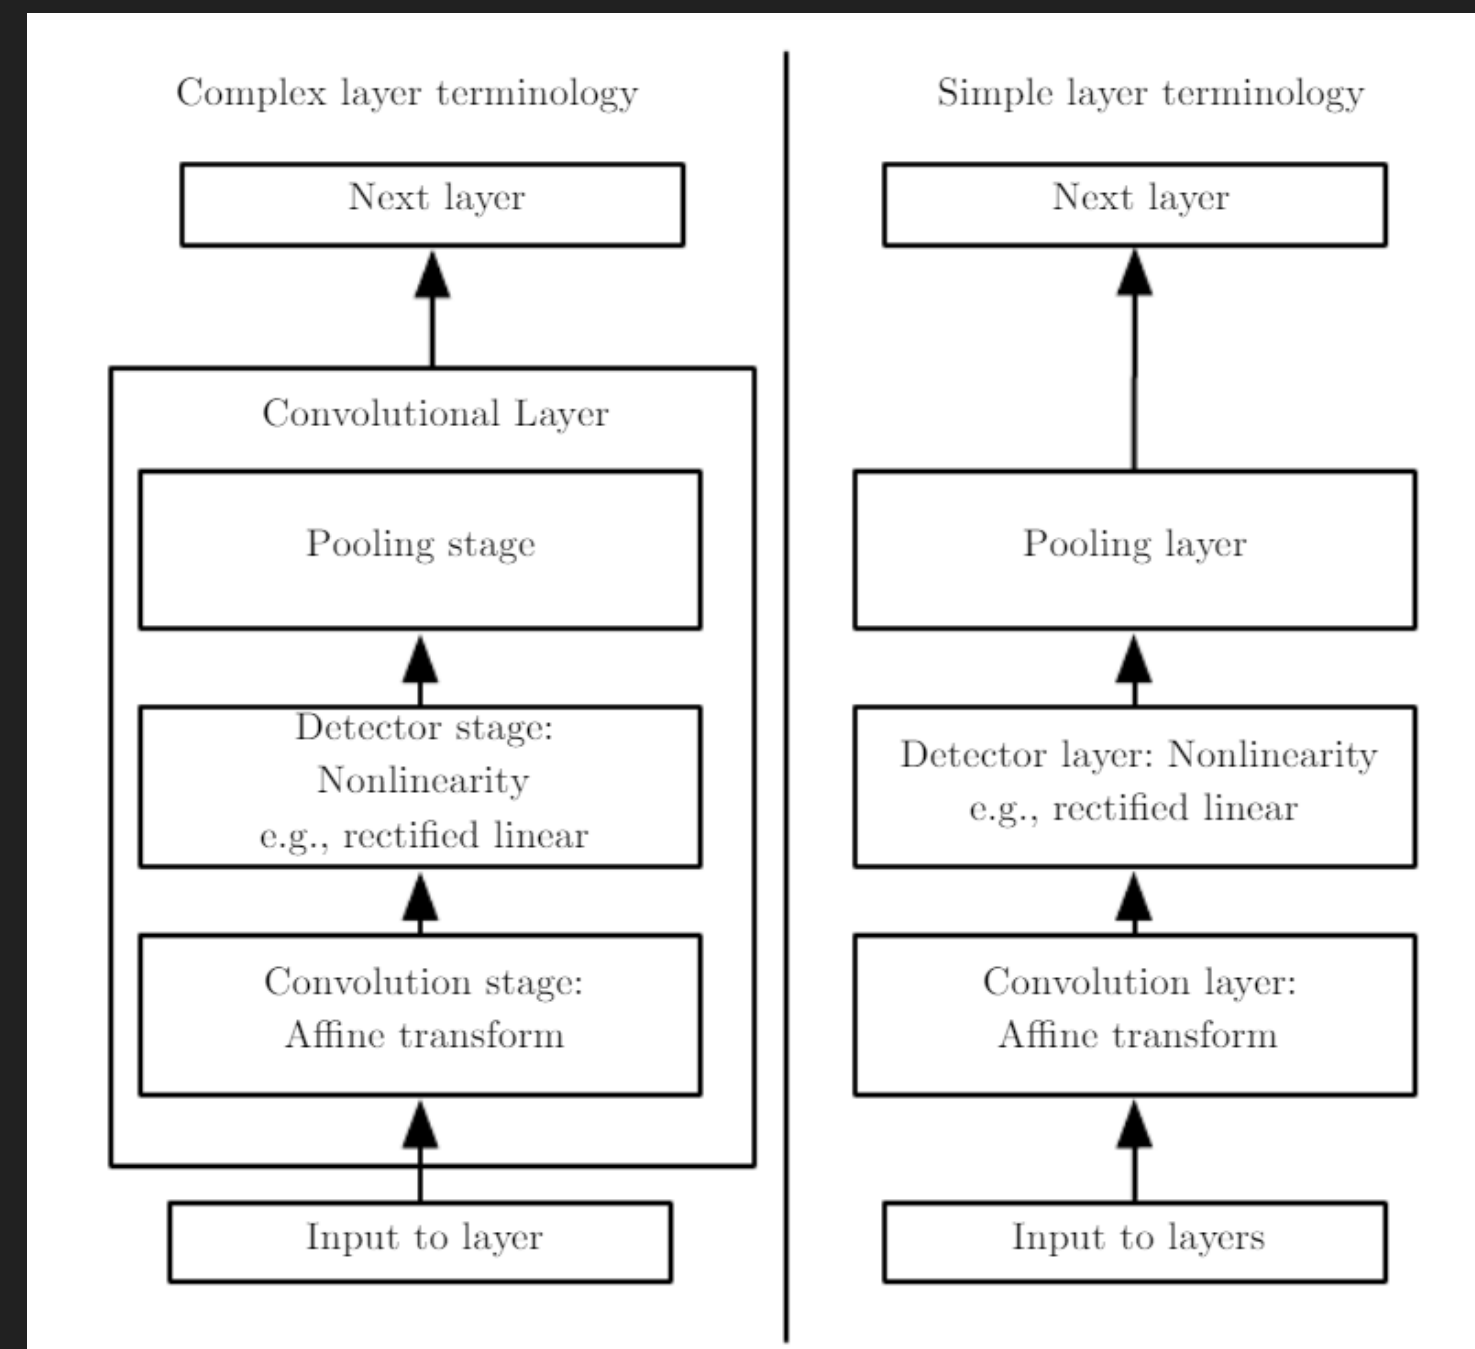
Inspired by **visual processing system** of the brain

## CNN AS FEATURE EXTRACTORS FOR DNNs



# LAYERS AND STRUCTURES

- Convolutional layer
- Pooling / subsampling layer








# OVERALL PREMISE

- Core building block and namesake of a CNN
- Number of customisable parameters in this type of layer that control a set of learnable **filters**
- Each filter has a **set of units** (neurons), which each apply the filter **kernel** to a specific area of the input data (its **local receptive field**)
- Filter are **convolved** with the input data to produce a **feature map**
- Each filter learns to activate when it sees a specific feature



# CONVOLUTIONS

An operation which describes the **mixing** of two functions or pieces of information.

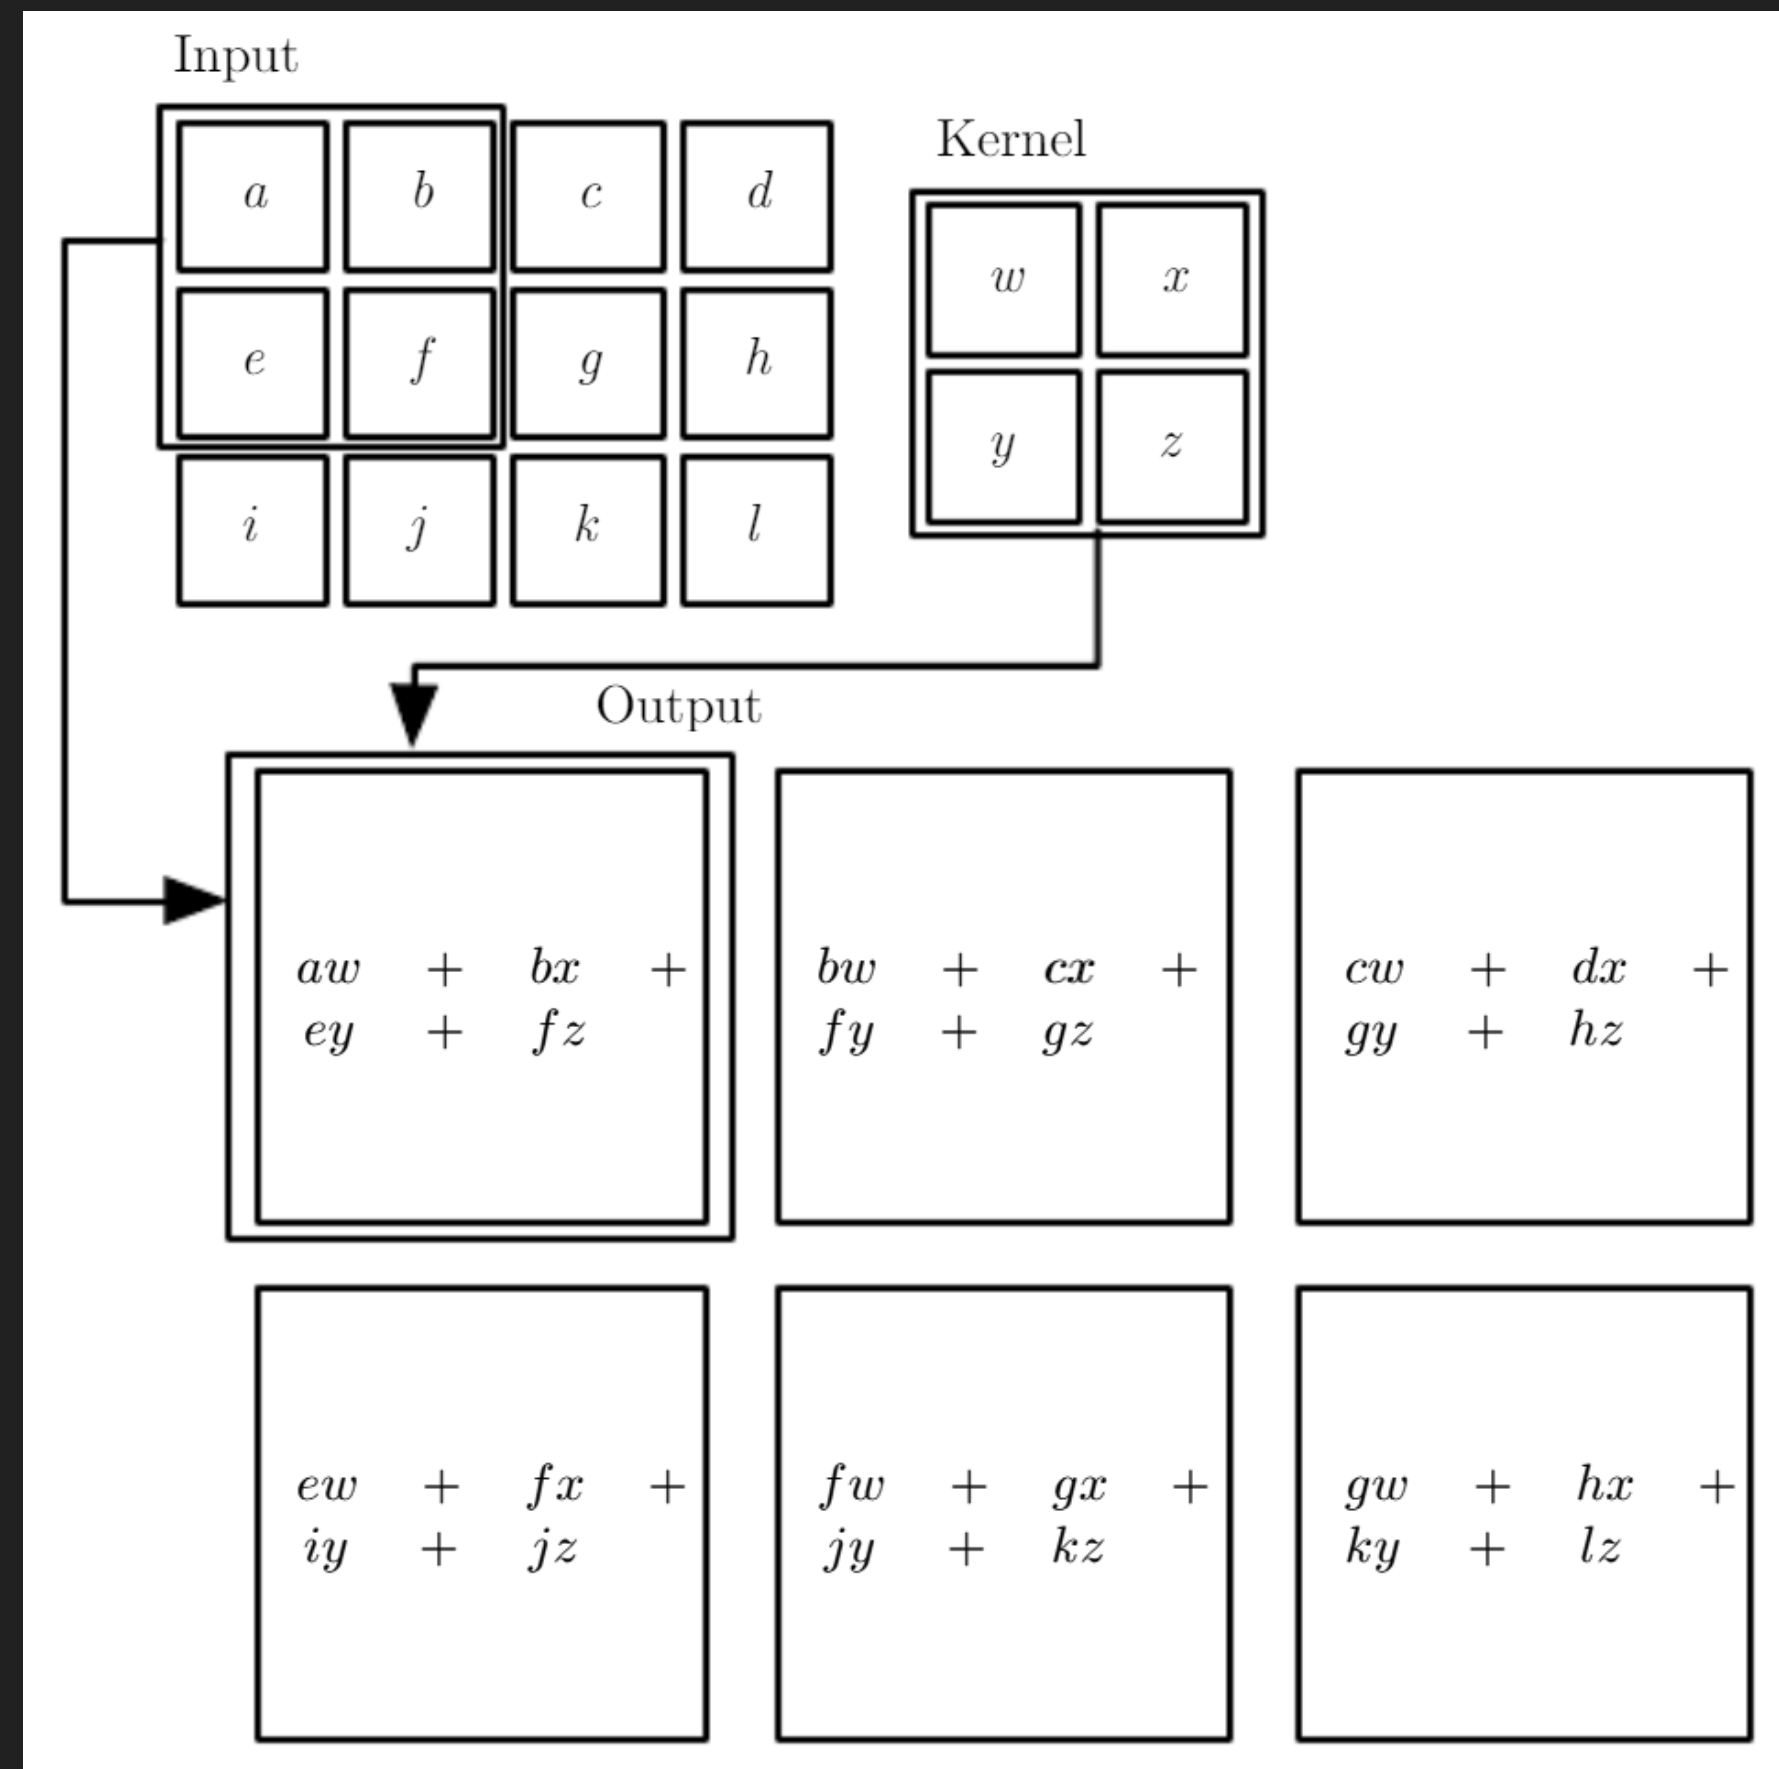
$$h(x) = f \otimes g = \int_{-\infty}^{\infty} f(x-u)g(u) du = \mathcal{F}^{-1} \left( \sqrt{2\pi} \mathcal{F}[f] \mathcal{F}[g] \right)$$

$$\text{feature map} = \text{input} \otimes \text{kernel} = \sum_{y=0}^{\text{columns}} \left( \sum_{x=0}^{\text{rows}} \text{input}(x-a, y-b) \text{kernel}(x, y) \right) = \mathcal{F}^{-1} \left( \sqrt{2\pi} \mathcal{F}[\text{input}] \mathcal{F}[\text{kernel}] \right)$$

Most libraries implement cross-correlation\*

\* other interpretations of convolutions can apply in other fields

# COMPUTING FEATURE MAPS

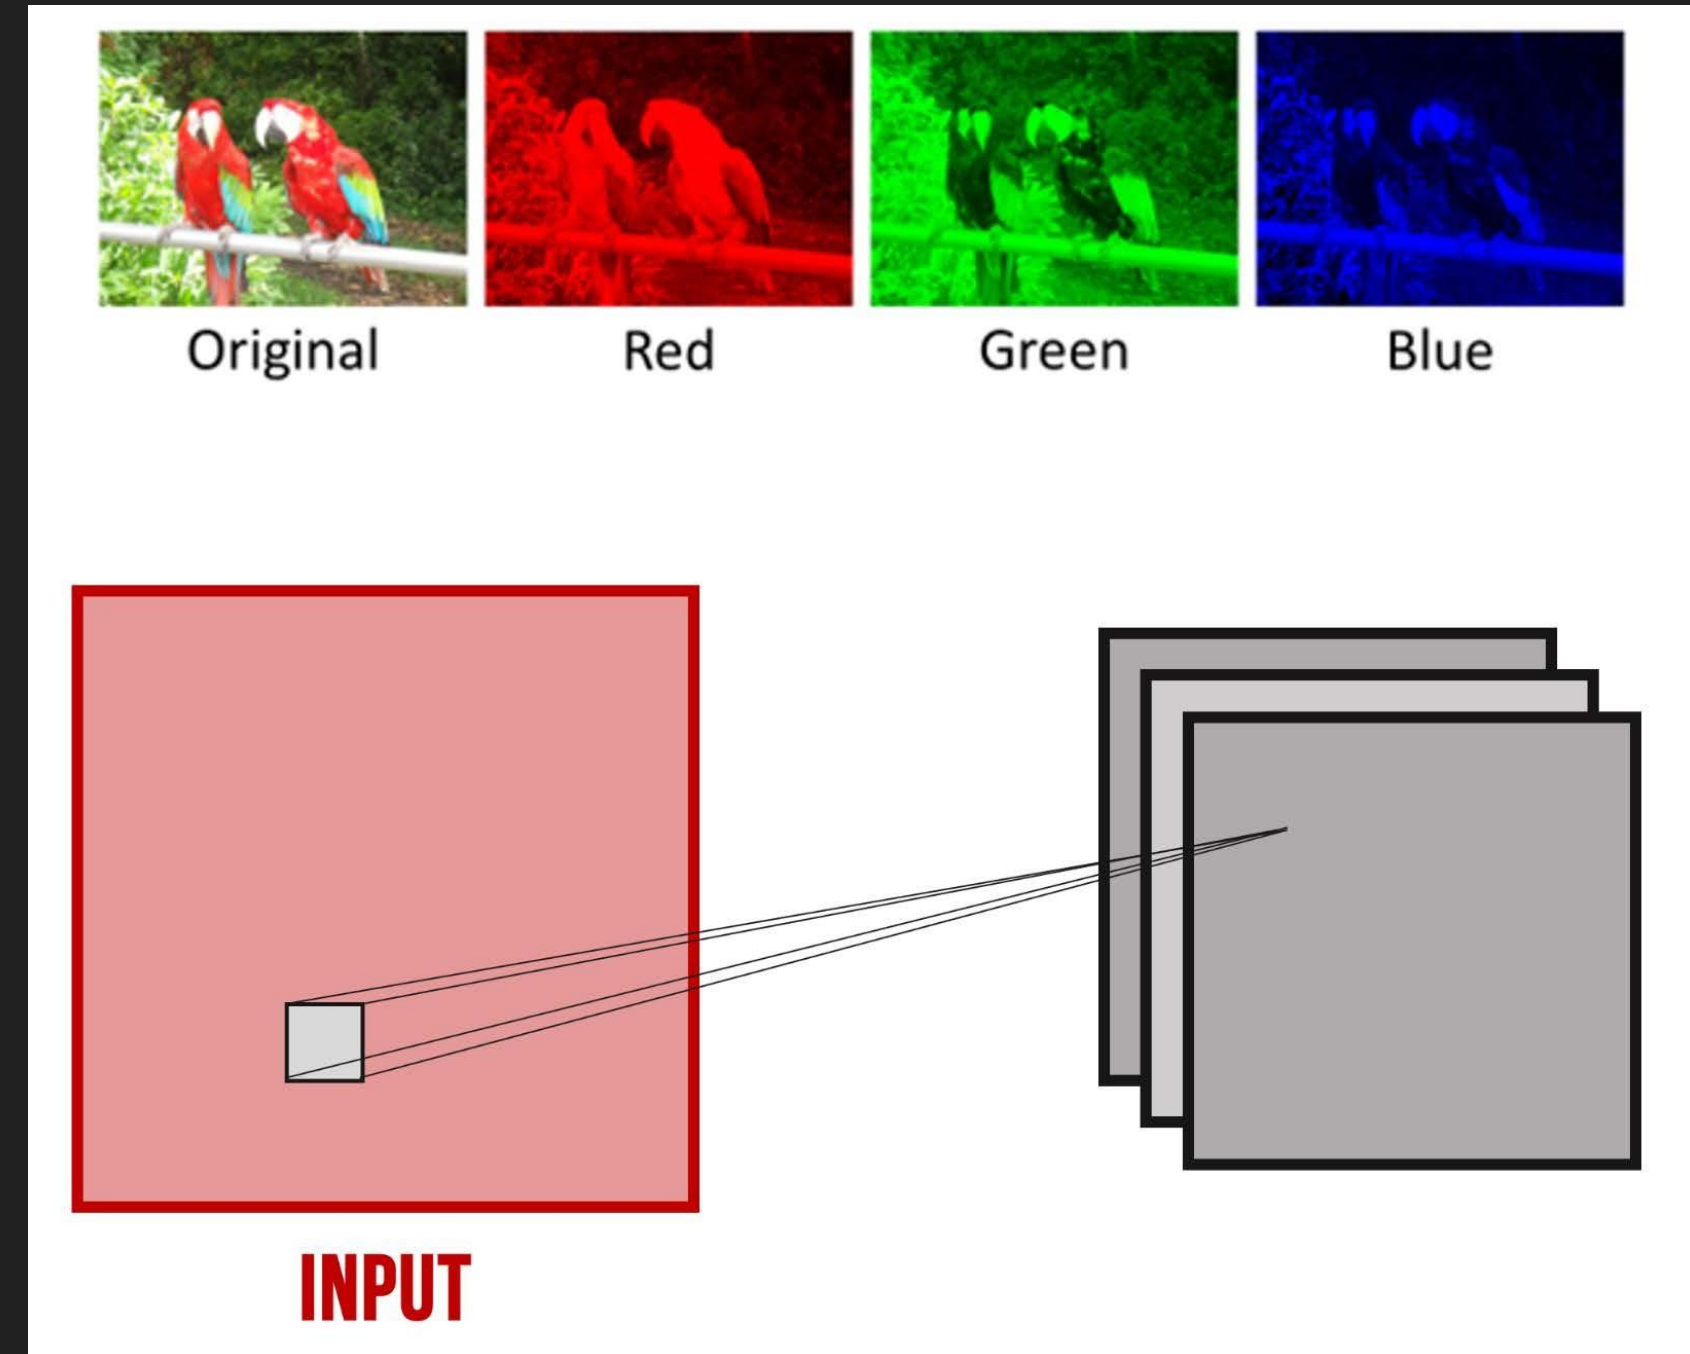


# INPUT

## Tensors

- 1D = sequences
- 2D = images
- 3D = video

Inputs can be multi-channel



## FILTERS

- Number of filters **equal** to the number of features you want the network to detect, you define how many you want
- Features are not defined; the CNN will learn them by learning the weights and biases of the filters

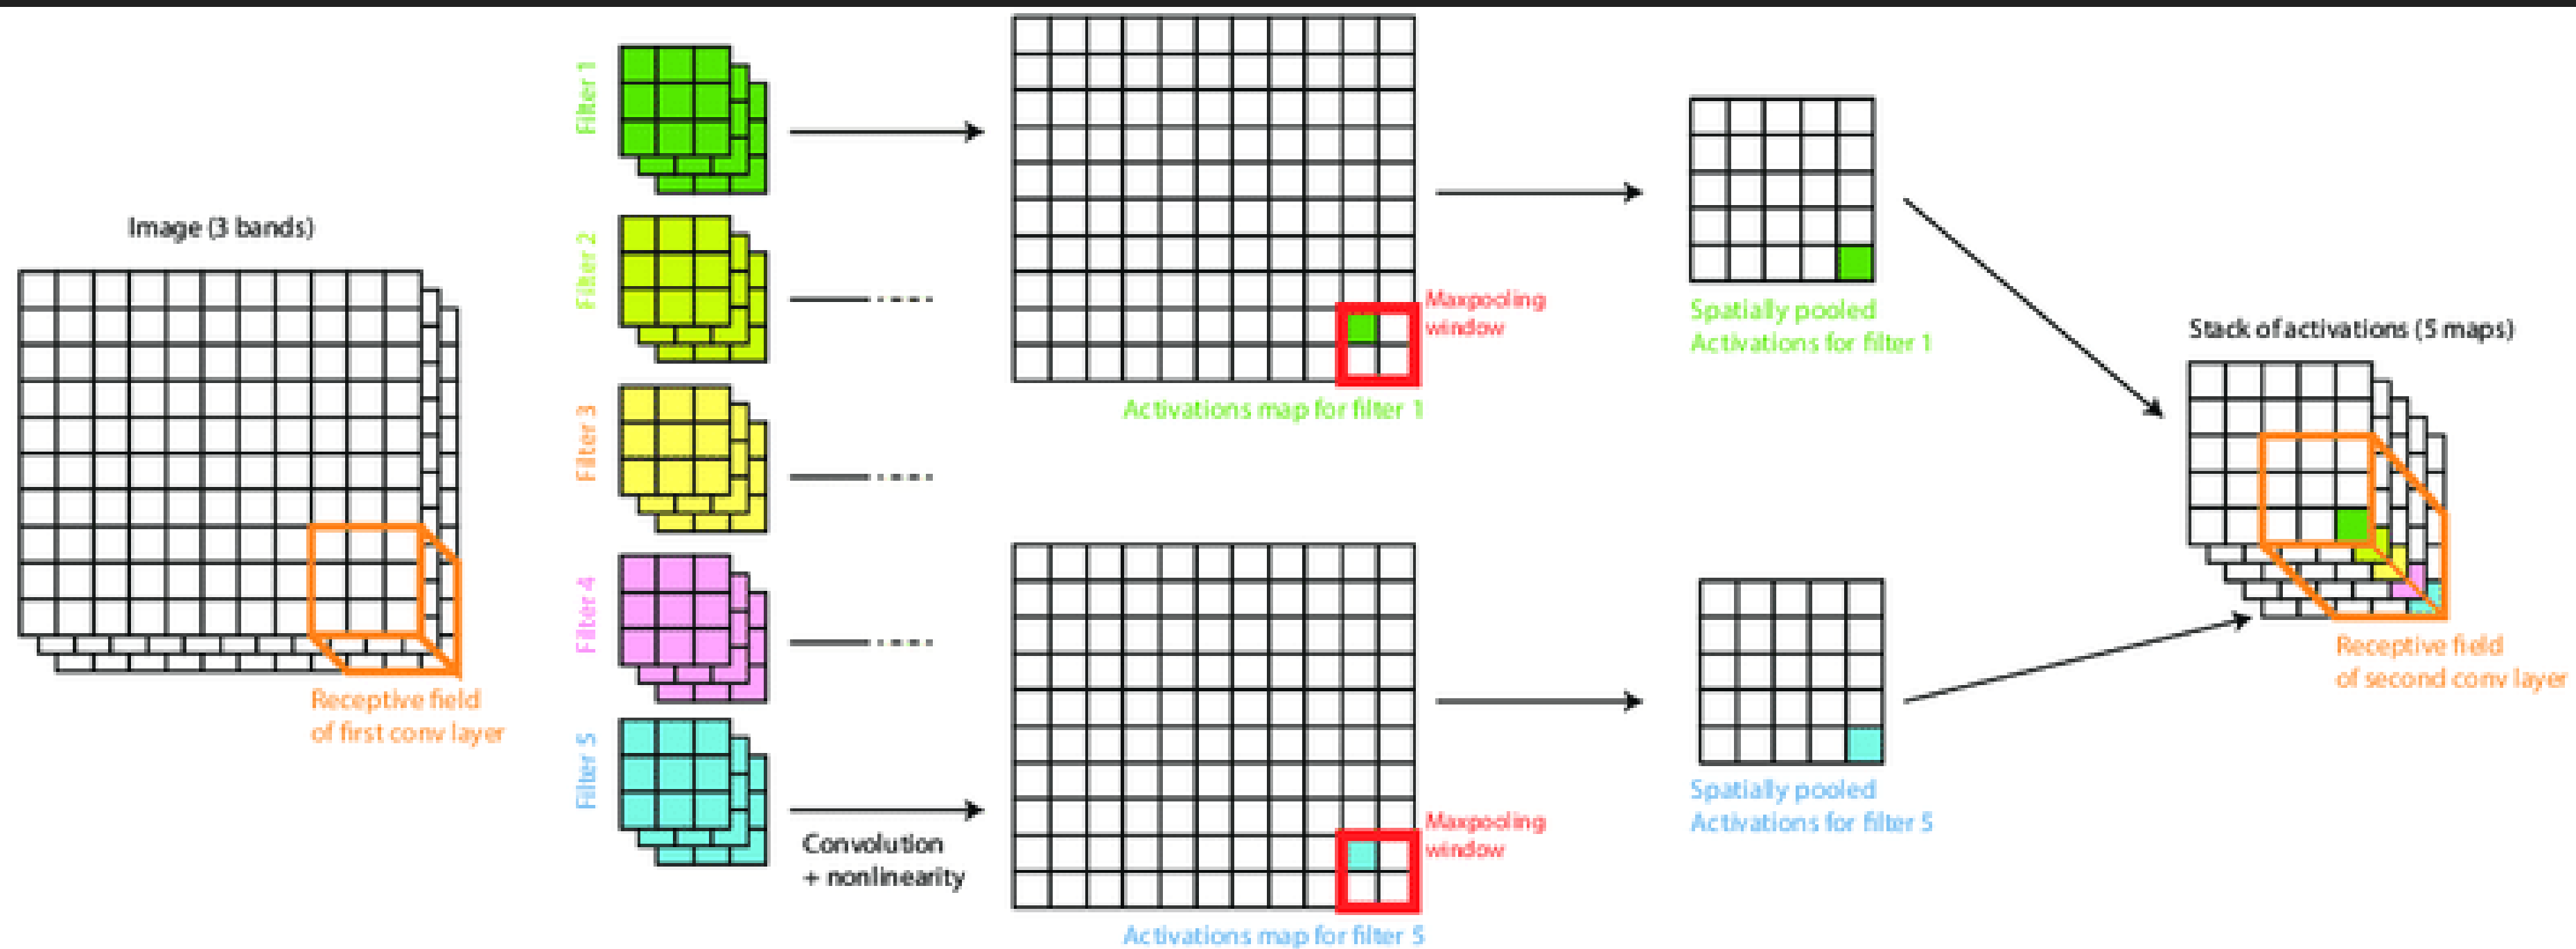
## NEURONS

- Filters convolve with the input to produce a series of neurons in an activation map that each detect a specific feature in a specific region of the input data
- Each neuron behaves the same by applying the filter's kernel, but on a different part of the input

## LOCAL RECEPTIVE FIELD


- The current part of the input that is being convolved with the filter



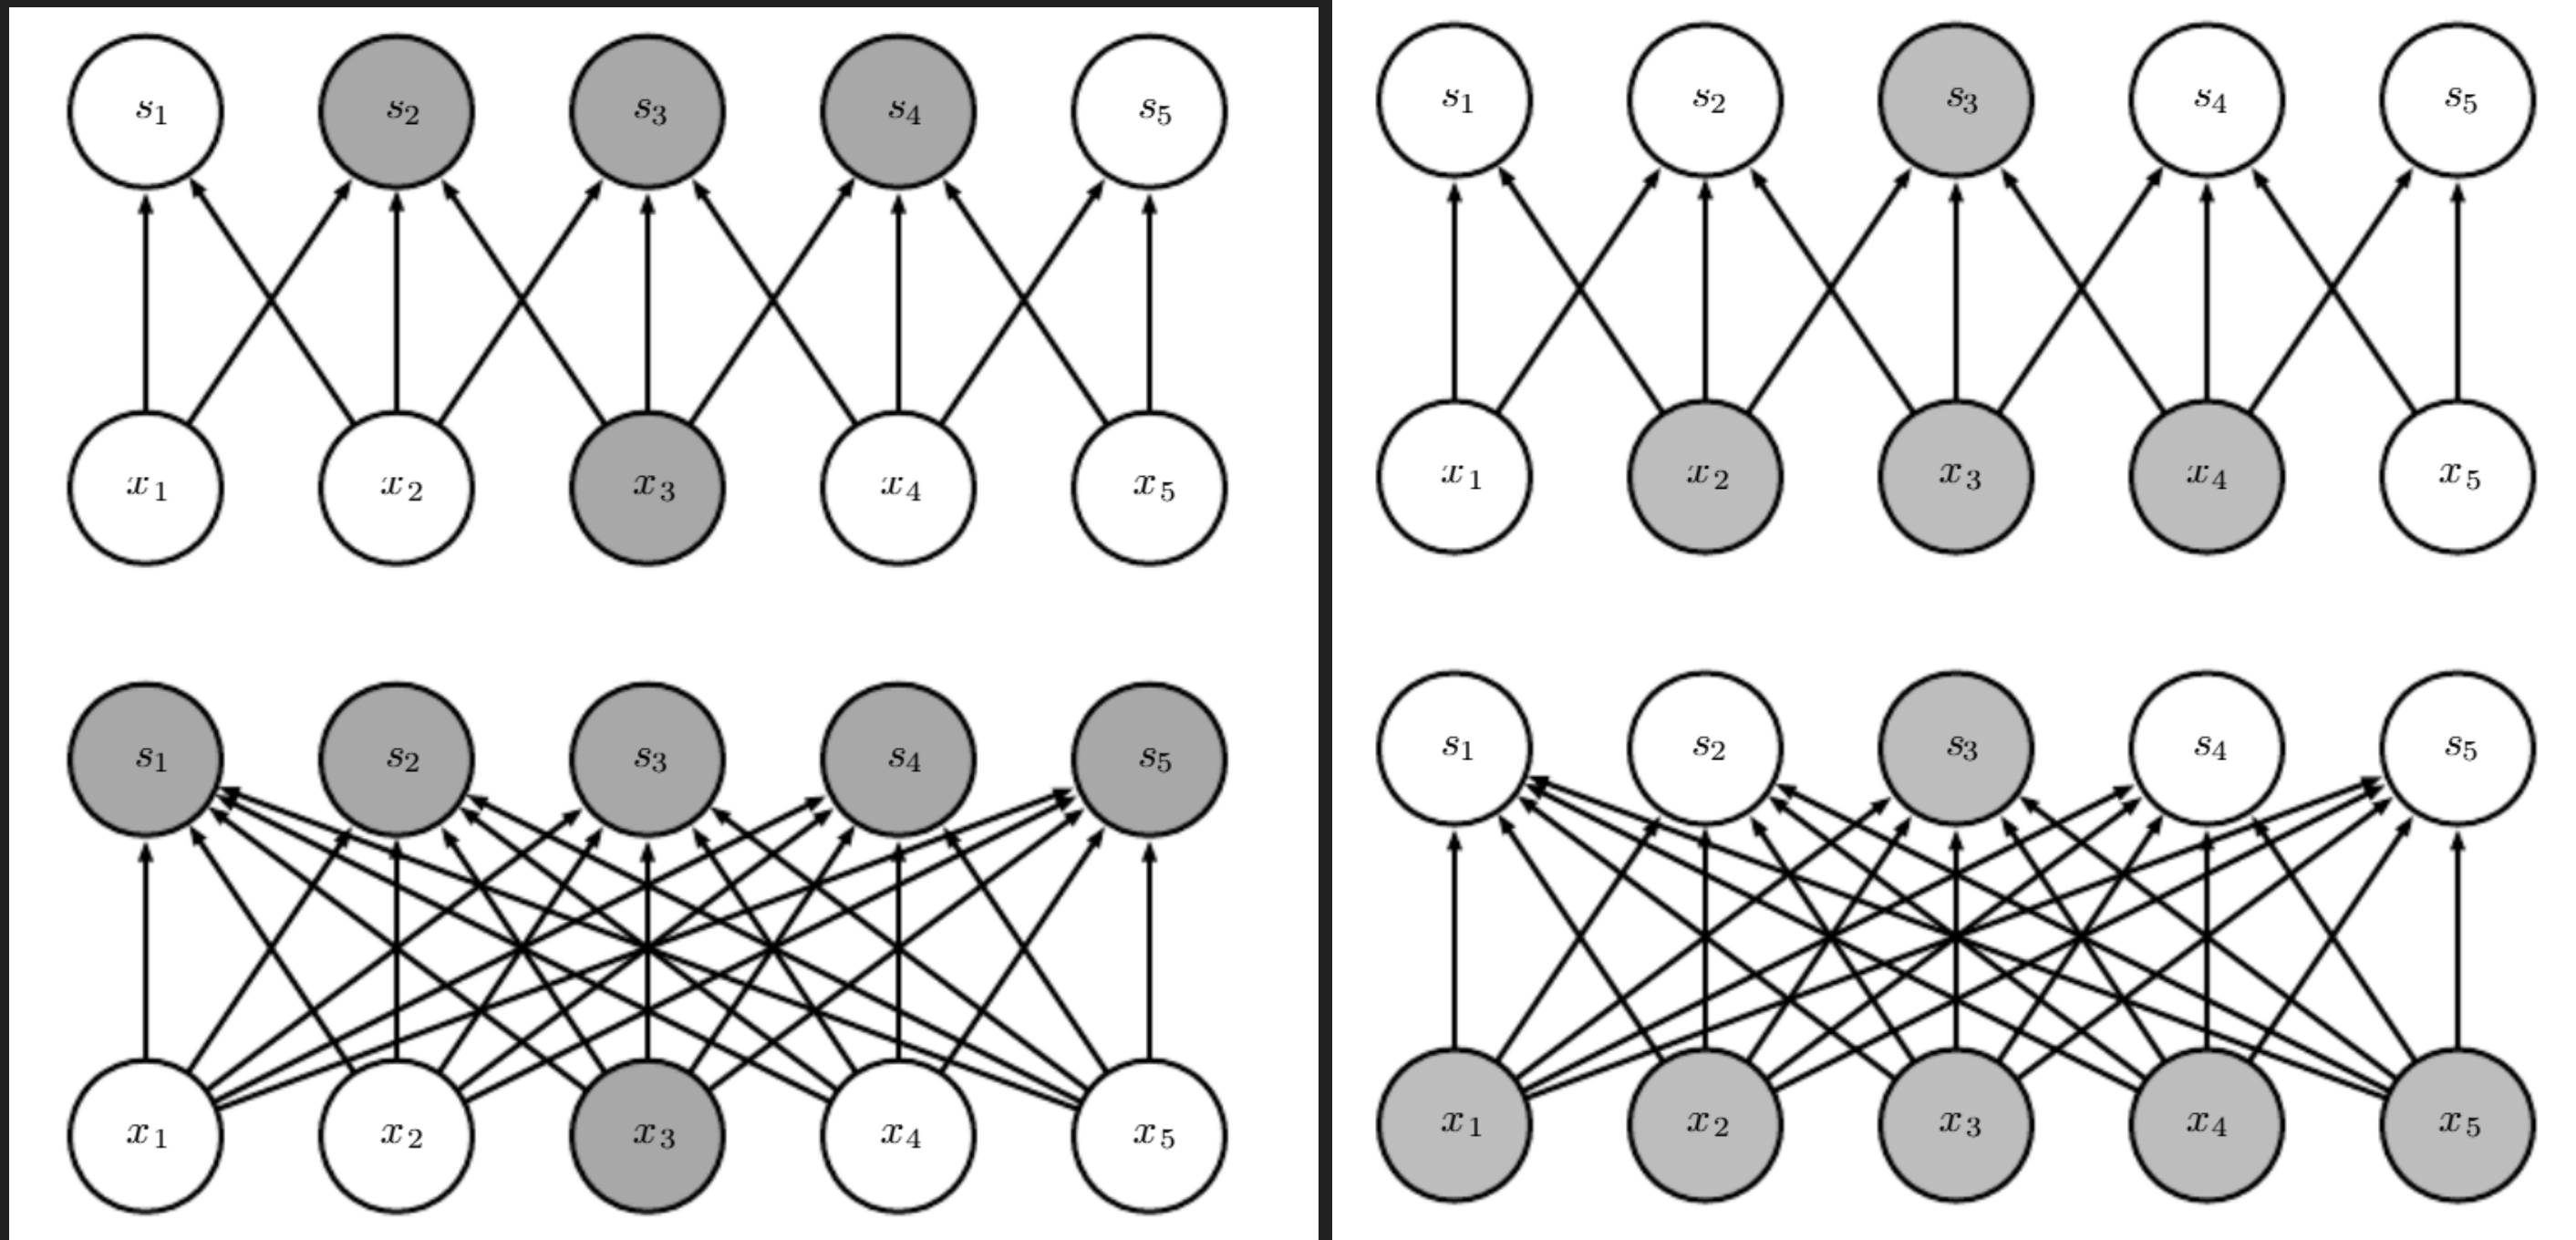


## PARAMETER SHARING



- Every neuron in a feature map has a unique local receptive field
- But, the weights and biases (filter) applied to these fields are shared
- This means that every unit in the filter detects exactly the same feature, just at different locations in the input



# SPARSITY



# HIERARCHY



# CONVOLUTIONS

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
1	0	1	1	0
0	1	1	0	0

5x5 input data

1	1	1
0	1	1
0	0	1

3x3 local receptive field

x1	x0	x1
x0	x1	x0
x1	x0	x1

Filter

4		

3x3 activation map



## COMPUTING FEATURE MAPS

1	1	1
0	1	1
0	0	1

Local Receptive Field  
3x3

x1	x0	x1
x0	x1	x0
x1	x0	x1

Filter  
3x3

$\Sigma$

1 × 1	1 × 0	1 × 1
0 × 0	1 × 1	1 × 0
0 × 1	0 × 0	1 × 1

=

4		

Activation Map

# CONVOLUTIONS

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0



Image

4		

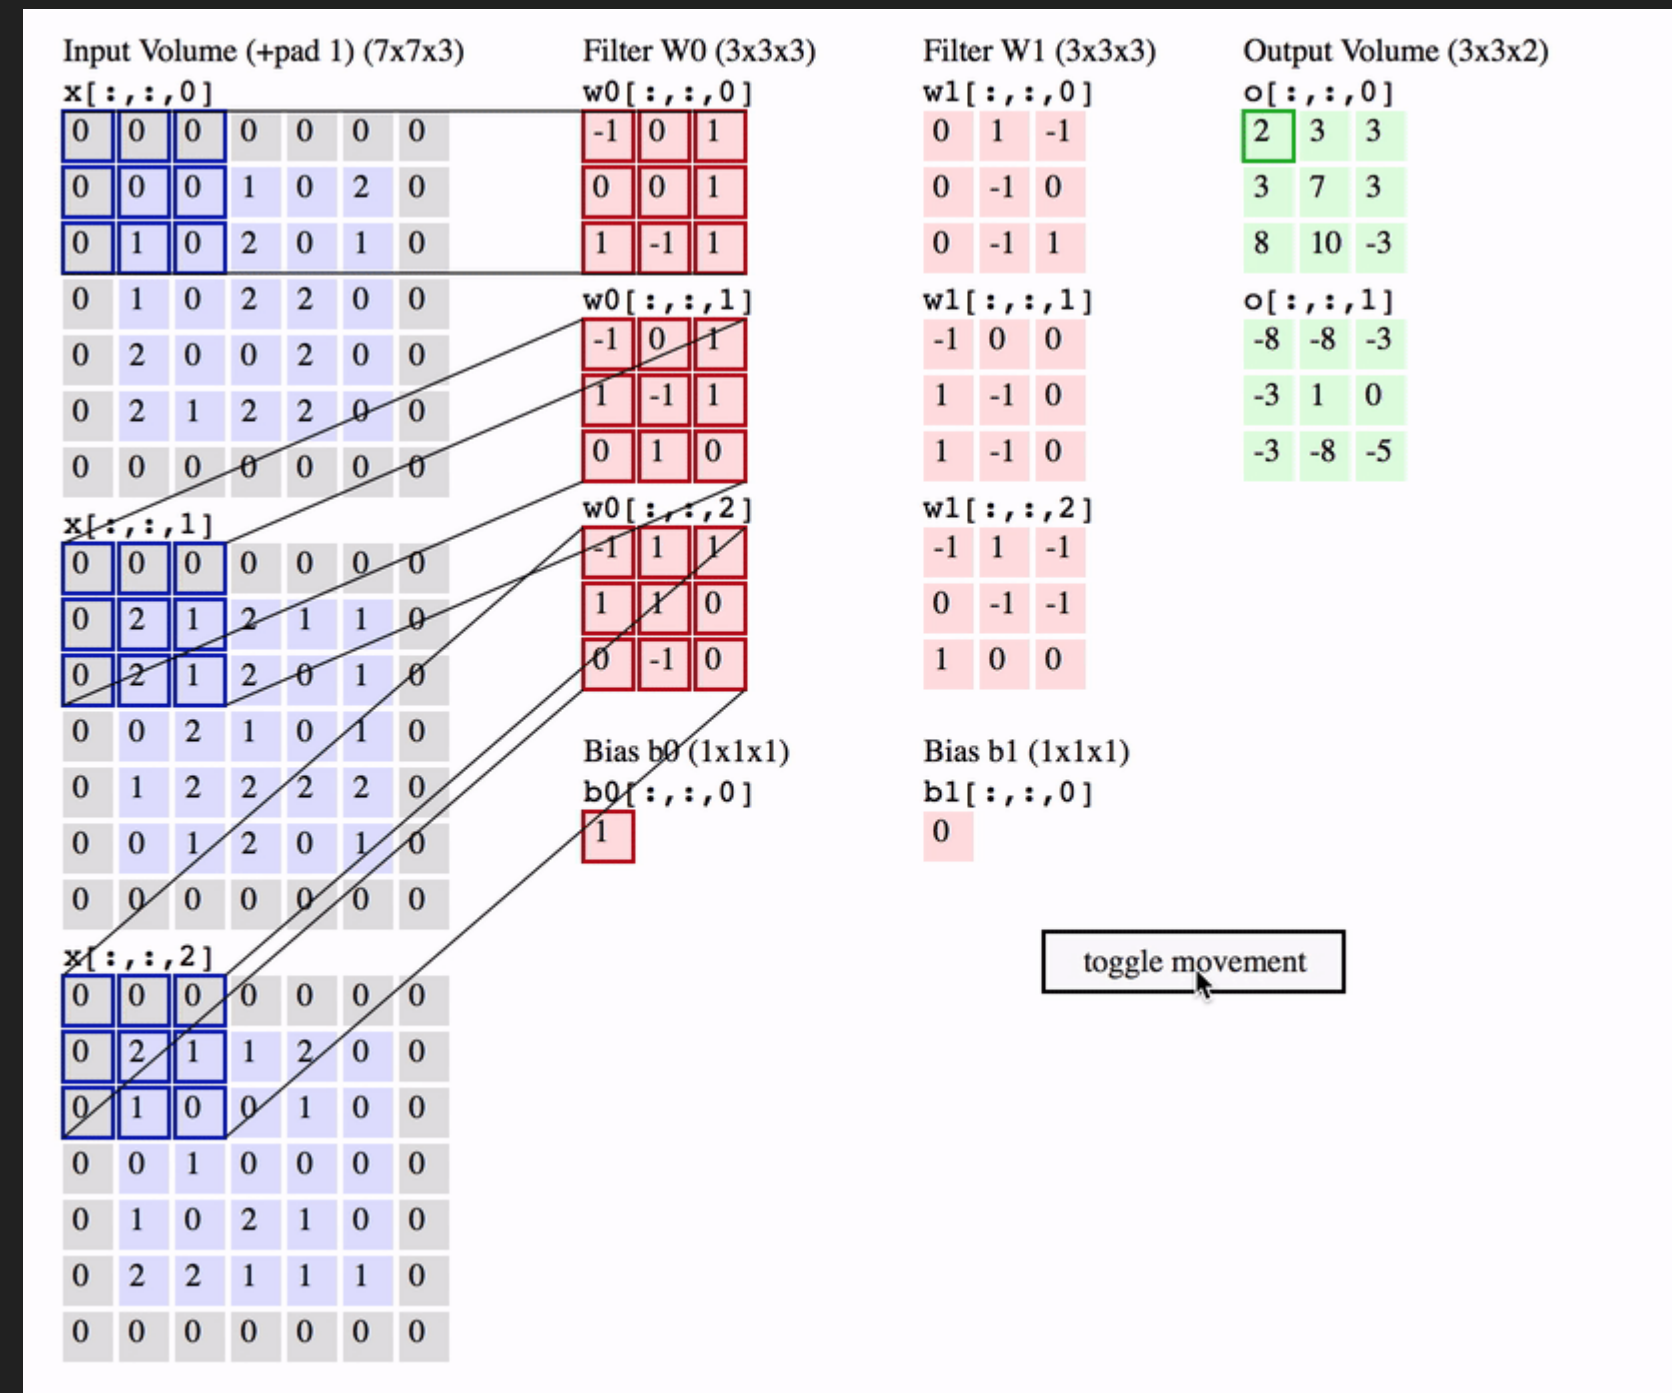
Convolved  
Feature

# STRIDES

- Size of receptive field matters; too small and it may not pick up a feature, too big and the network becomes fully-connected
- Width and height customisable; depth usually equal to input channels
- Stride length controls how far local receptive field slides; smaller number increases overlap between units in convolution layer

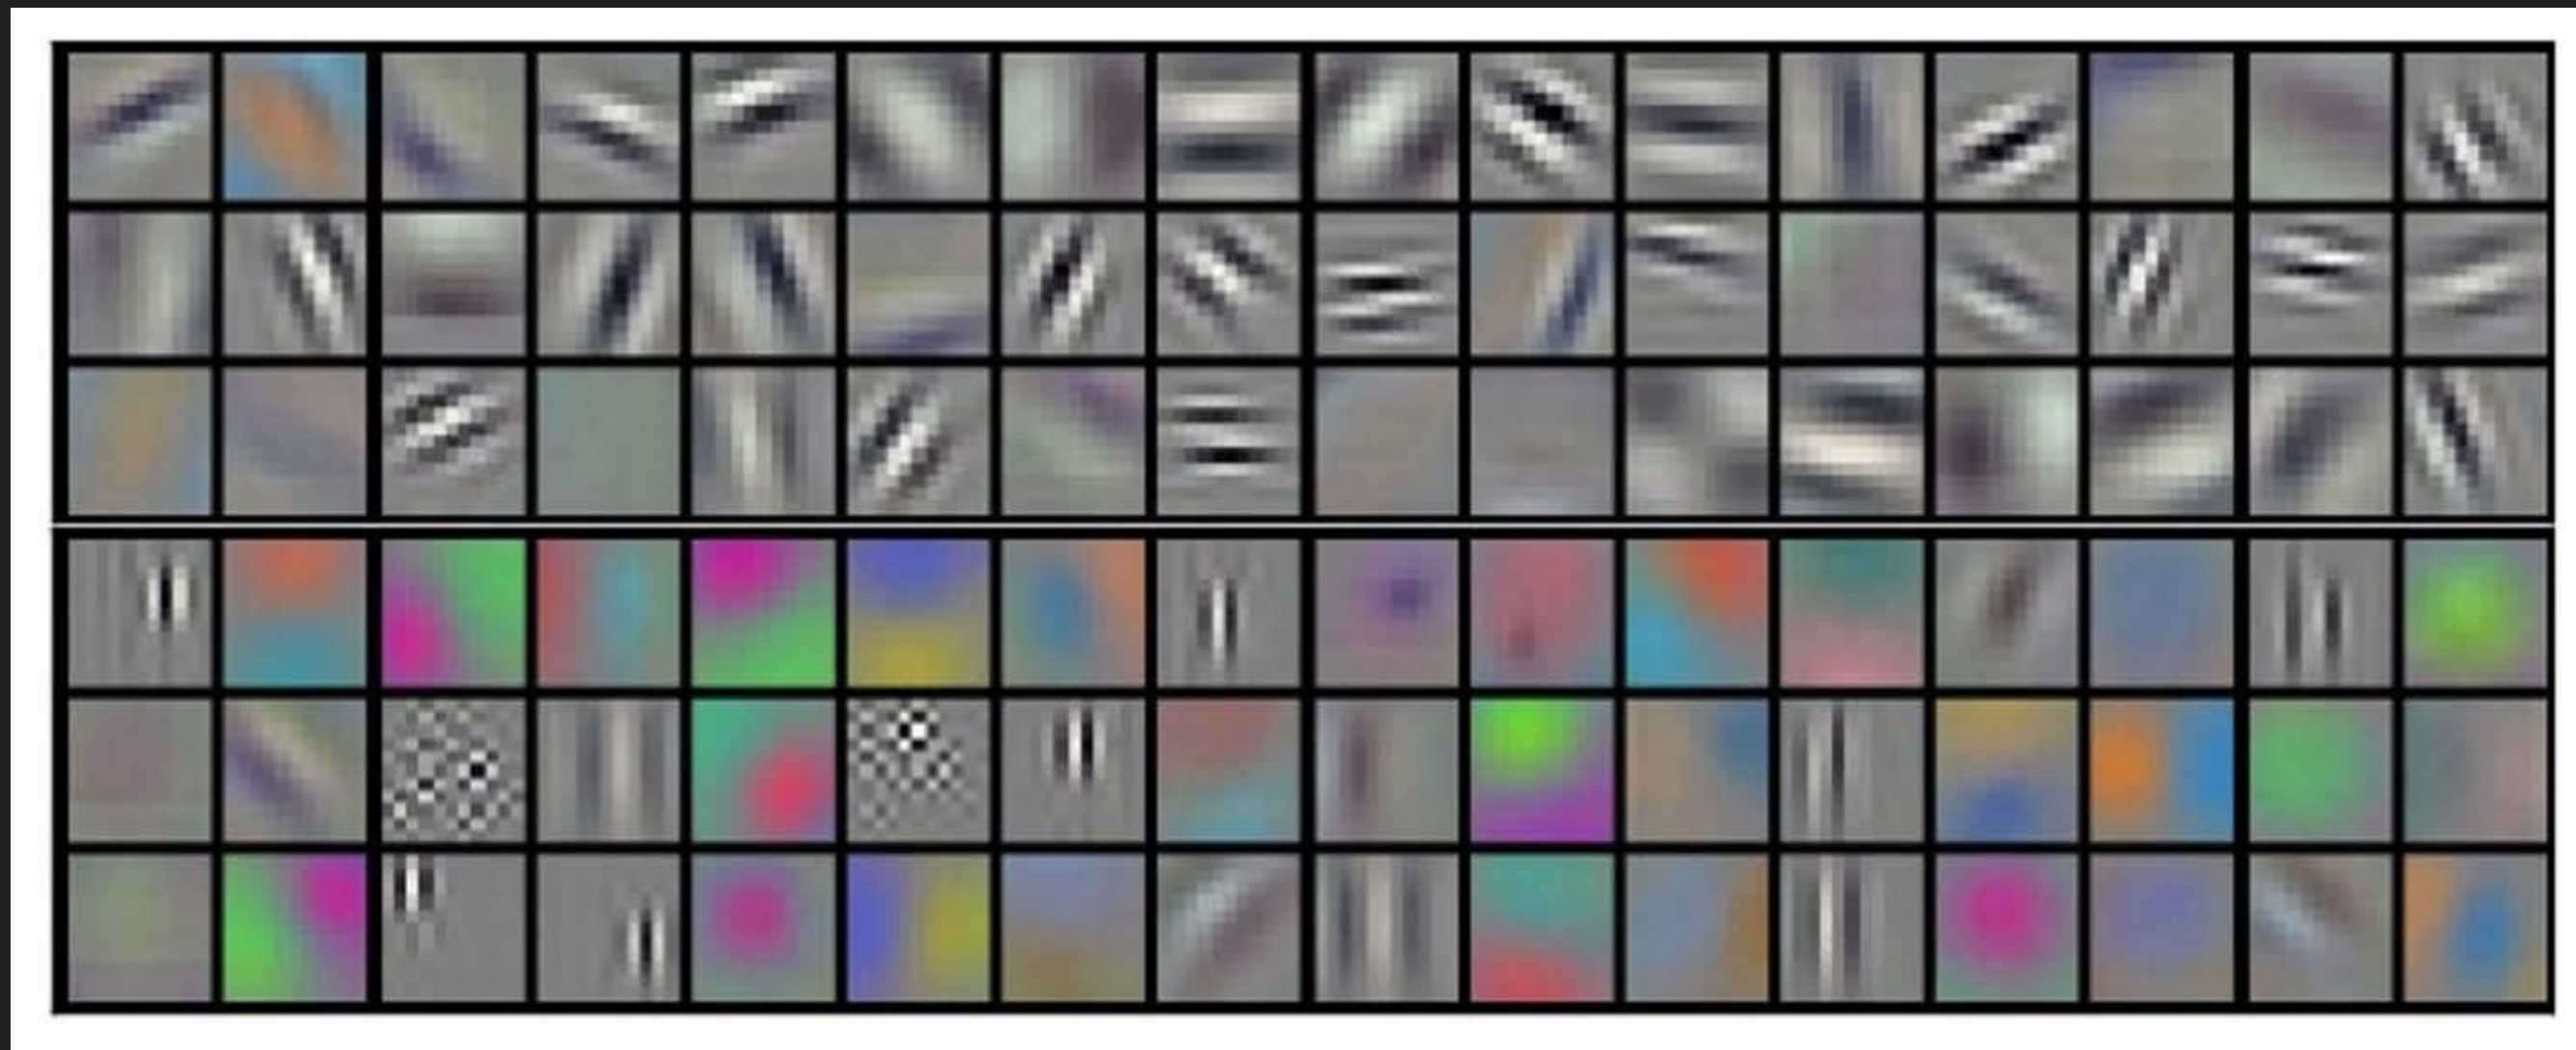


# MULTI-CHANNEL CONVOLUTIONS



## FEATURE MAPS

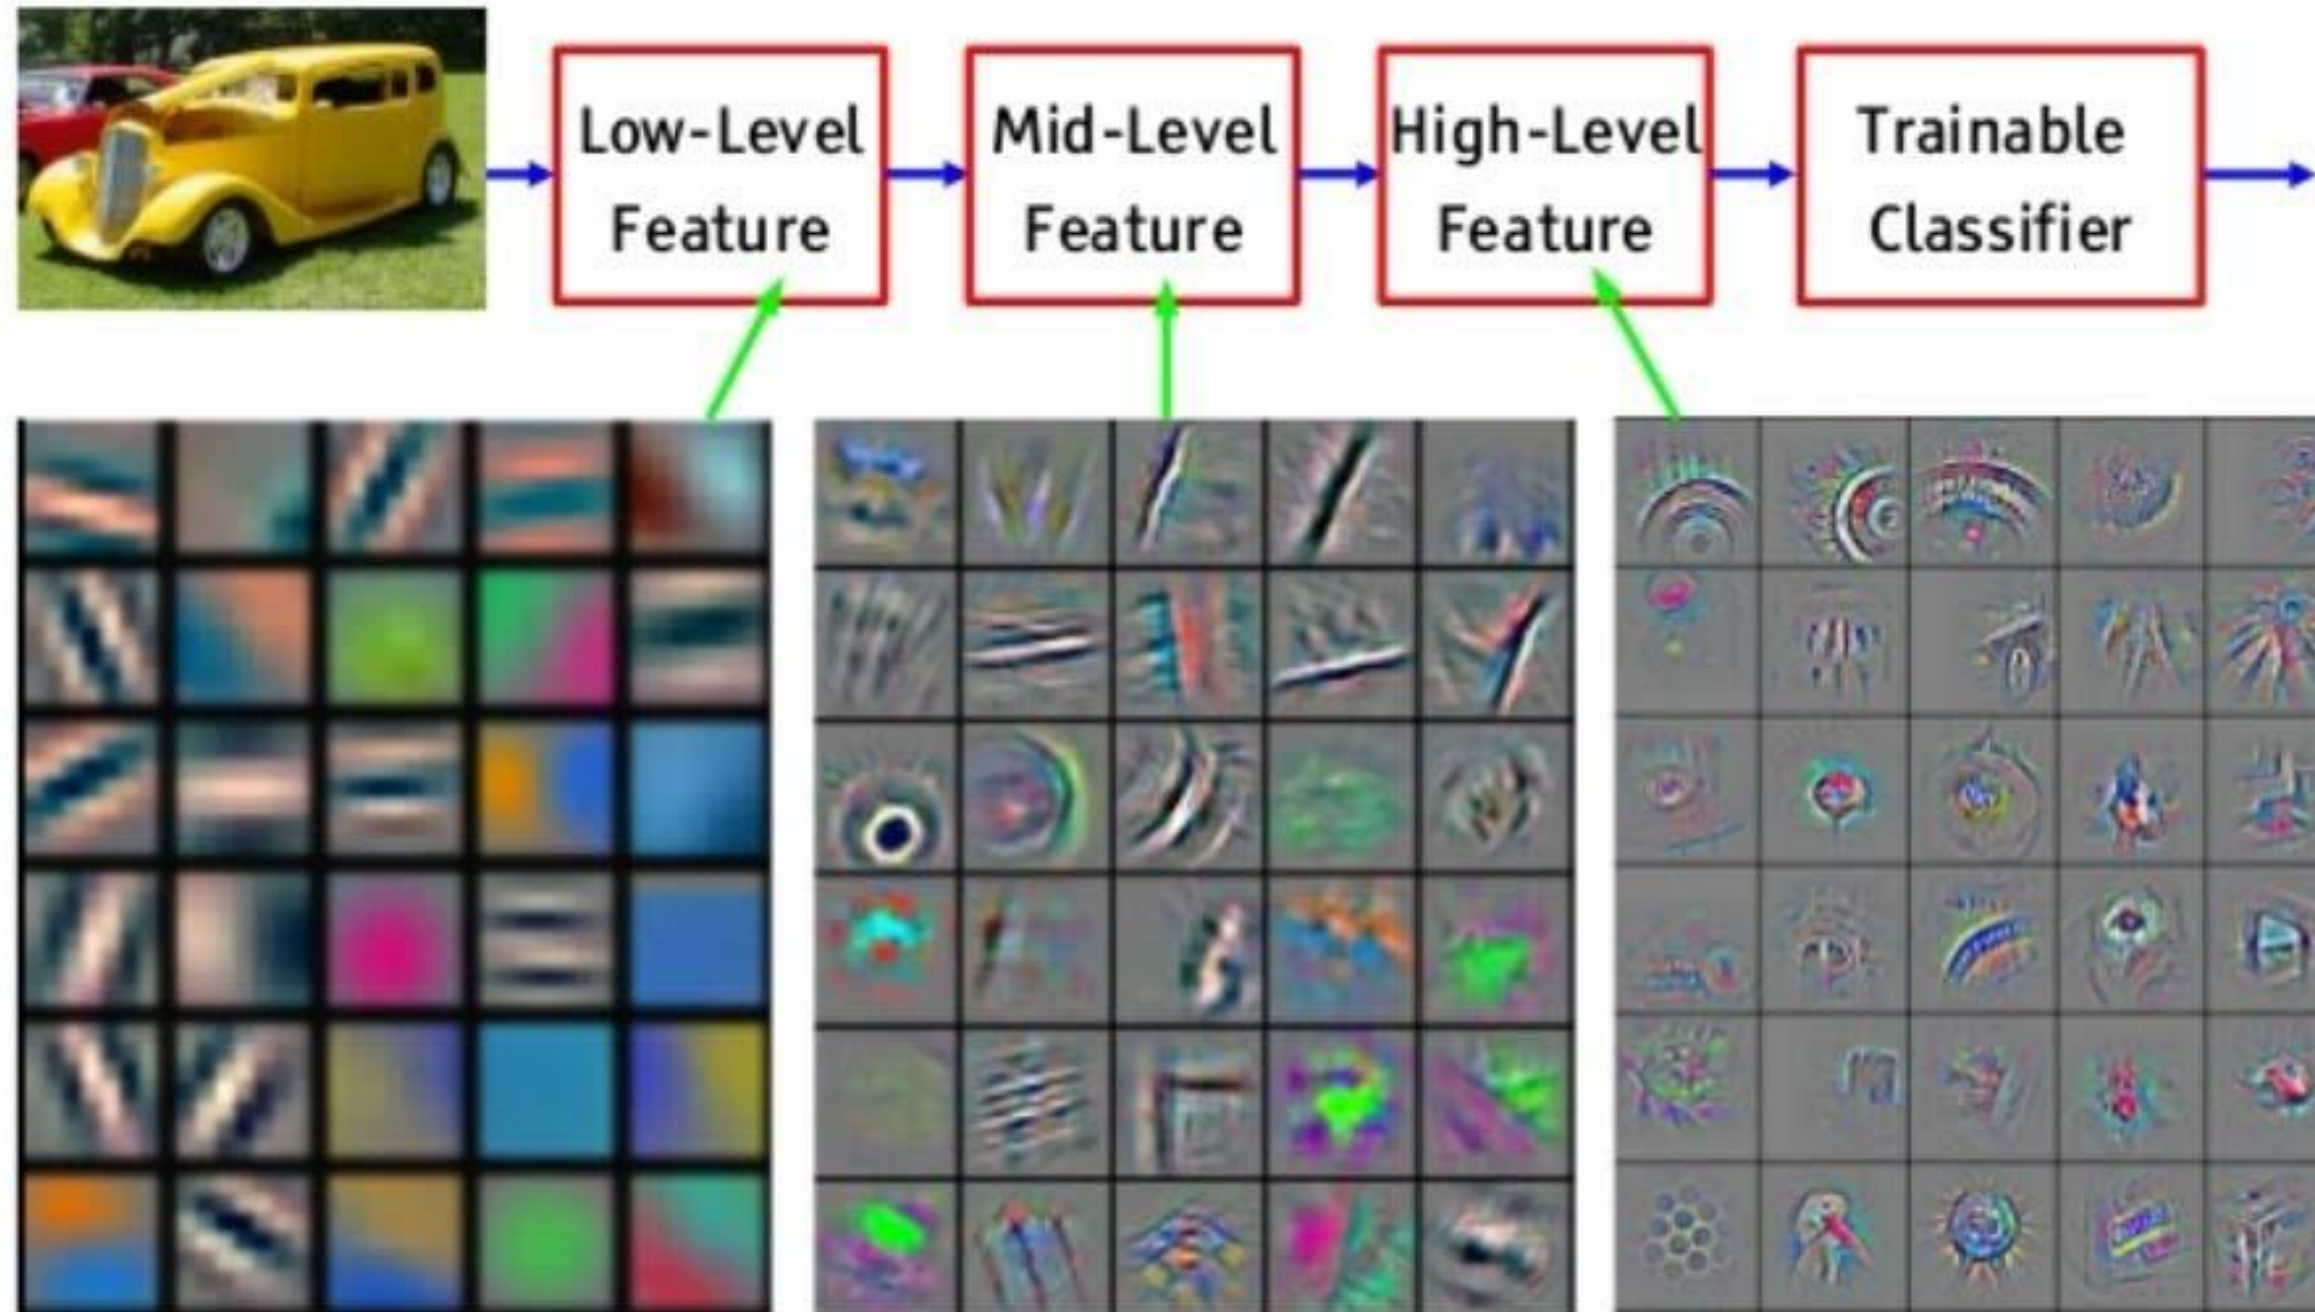
- Output of convolutions
- Number of feature maps produced equals number of features wanted to be detected (number of filters)
- Represents presence of feature at a given neuron
- Serves as input to the next layer (typically a pooling layer)





# HIERARCHICAL FEATURE LEARNING

■ It's **deep** if it has **more than one stage** of non-linear feature transformation



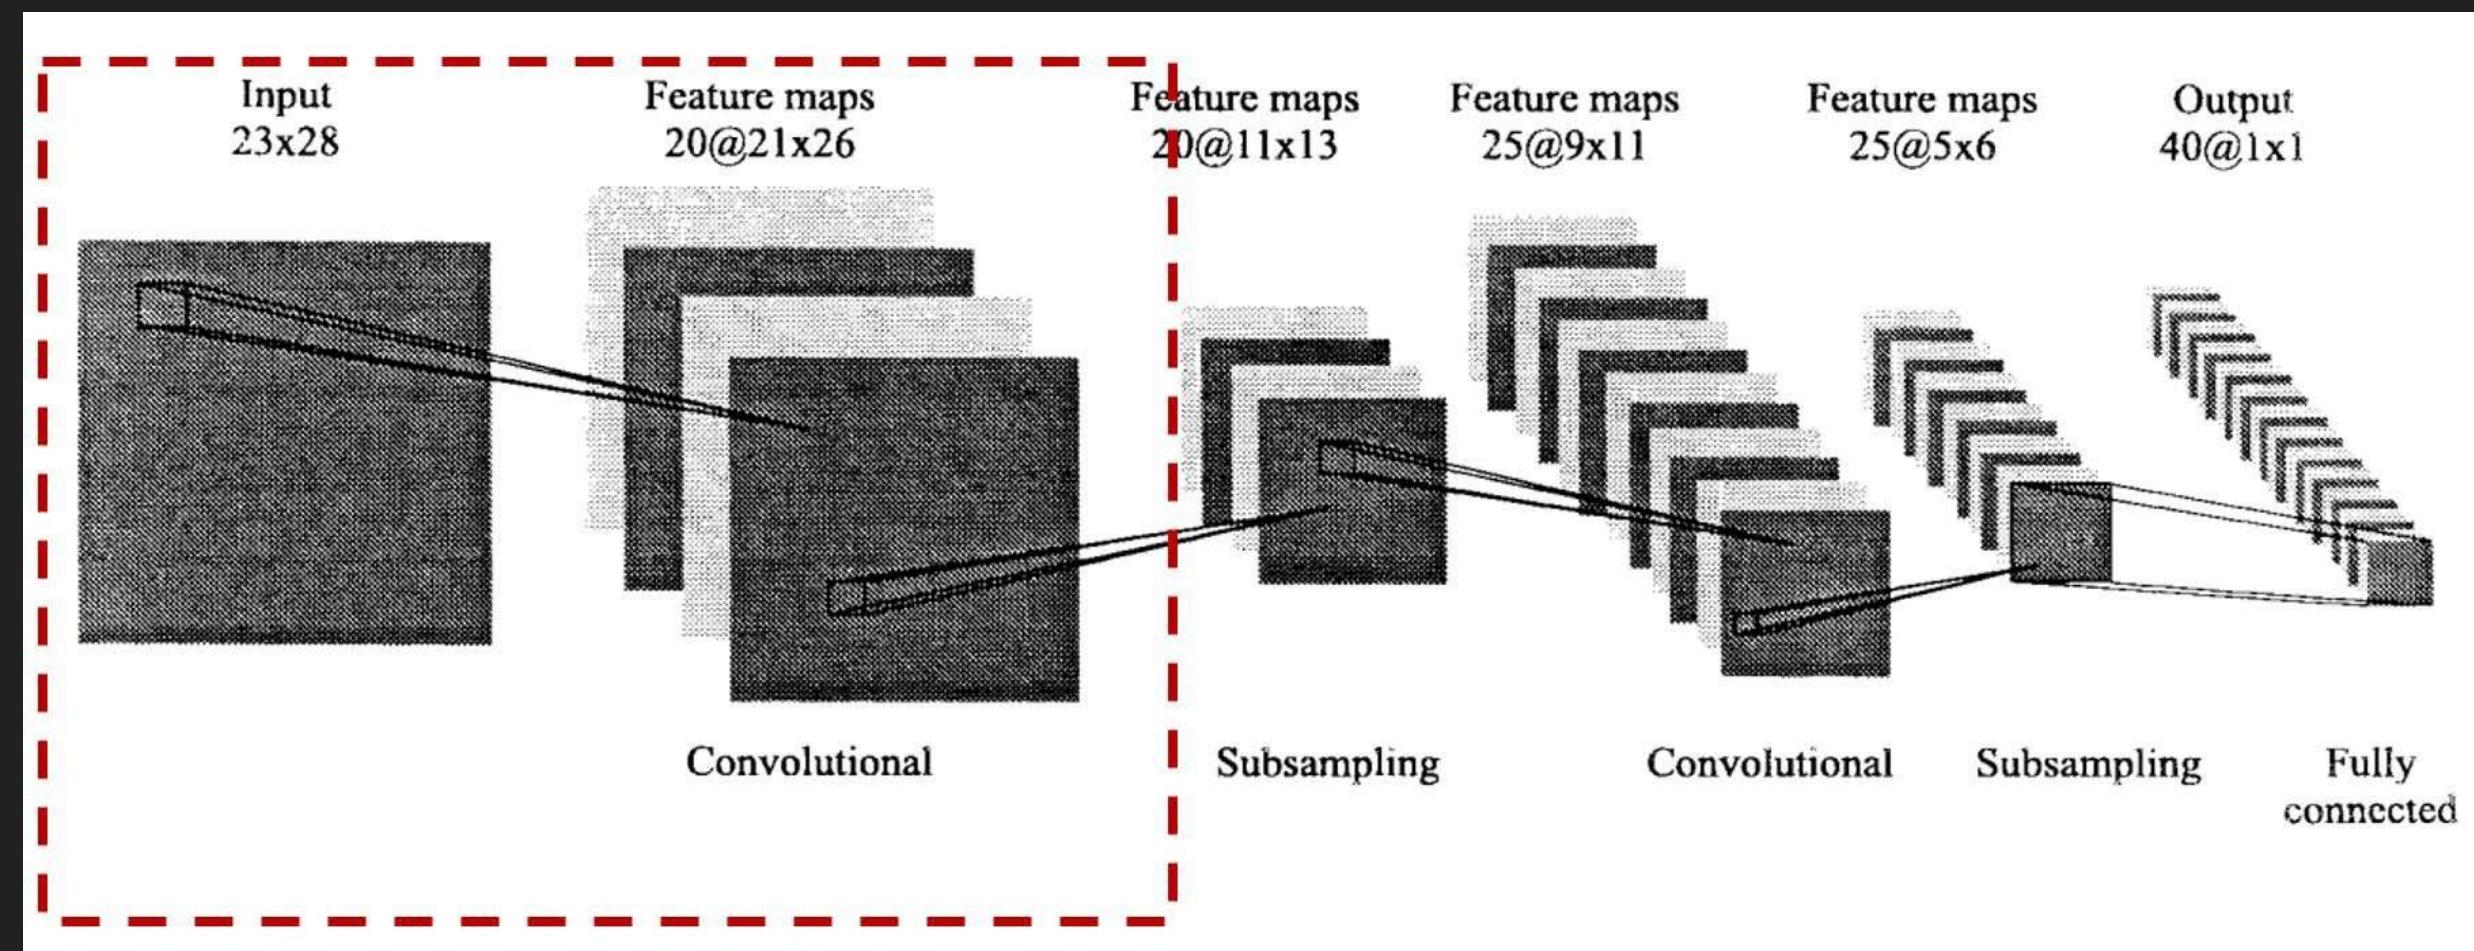
Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

## EQUIVARIANCE

- Translation equivariance
- Feature detected in one part of the data can appear in any other part of the data
- Creates a feature detector from one small sample that can be re-used elsewhere in the data
- Take the learned features and ‘convolve’ them with the larger dataset, getting different feature activation value at each location




# LAYER





## OVERALL PREMISE

- Produces a **summary statistic** of the output from the previous convolution layer
- Helps learned representation to be invariant to small translations of input
- Can **downsample** feature map into a condensed version
- Assists in controlling overfitting; reduces network parameters, size and computational cost



# POOLING

Configurable pooling window and stride

- max-pooling
- L2-pooling
- average-pooling
- overlapping pooling

3	4	4	0
2	3	3	1
0	0	1	1
3	1	0	2

**ACTIVATION MAP**



3	2
1	1

**POOLED  
ACTIVATION MAP**



## MAX POOLING

3	4	4	0
2	3	3	1
0	0	1	1
3	1	0	2

$$x_1 = \max\{3, 4, 2, 3\}$$

$$x_2 = \max\{4, 0, 3, 1\}$$


$$x_3 = \max\{0, 0, 3, 1\}$$

$$x_4 = \max\{1, 1, 0, 2\}$$

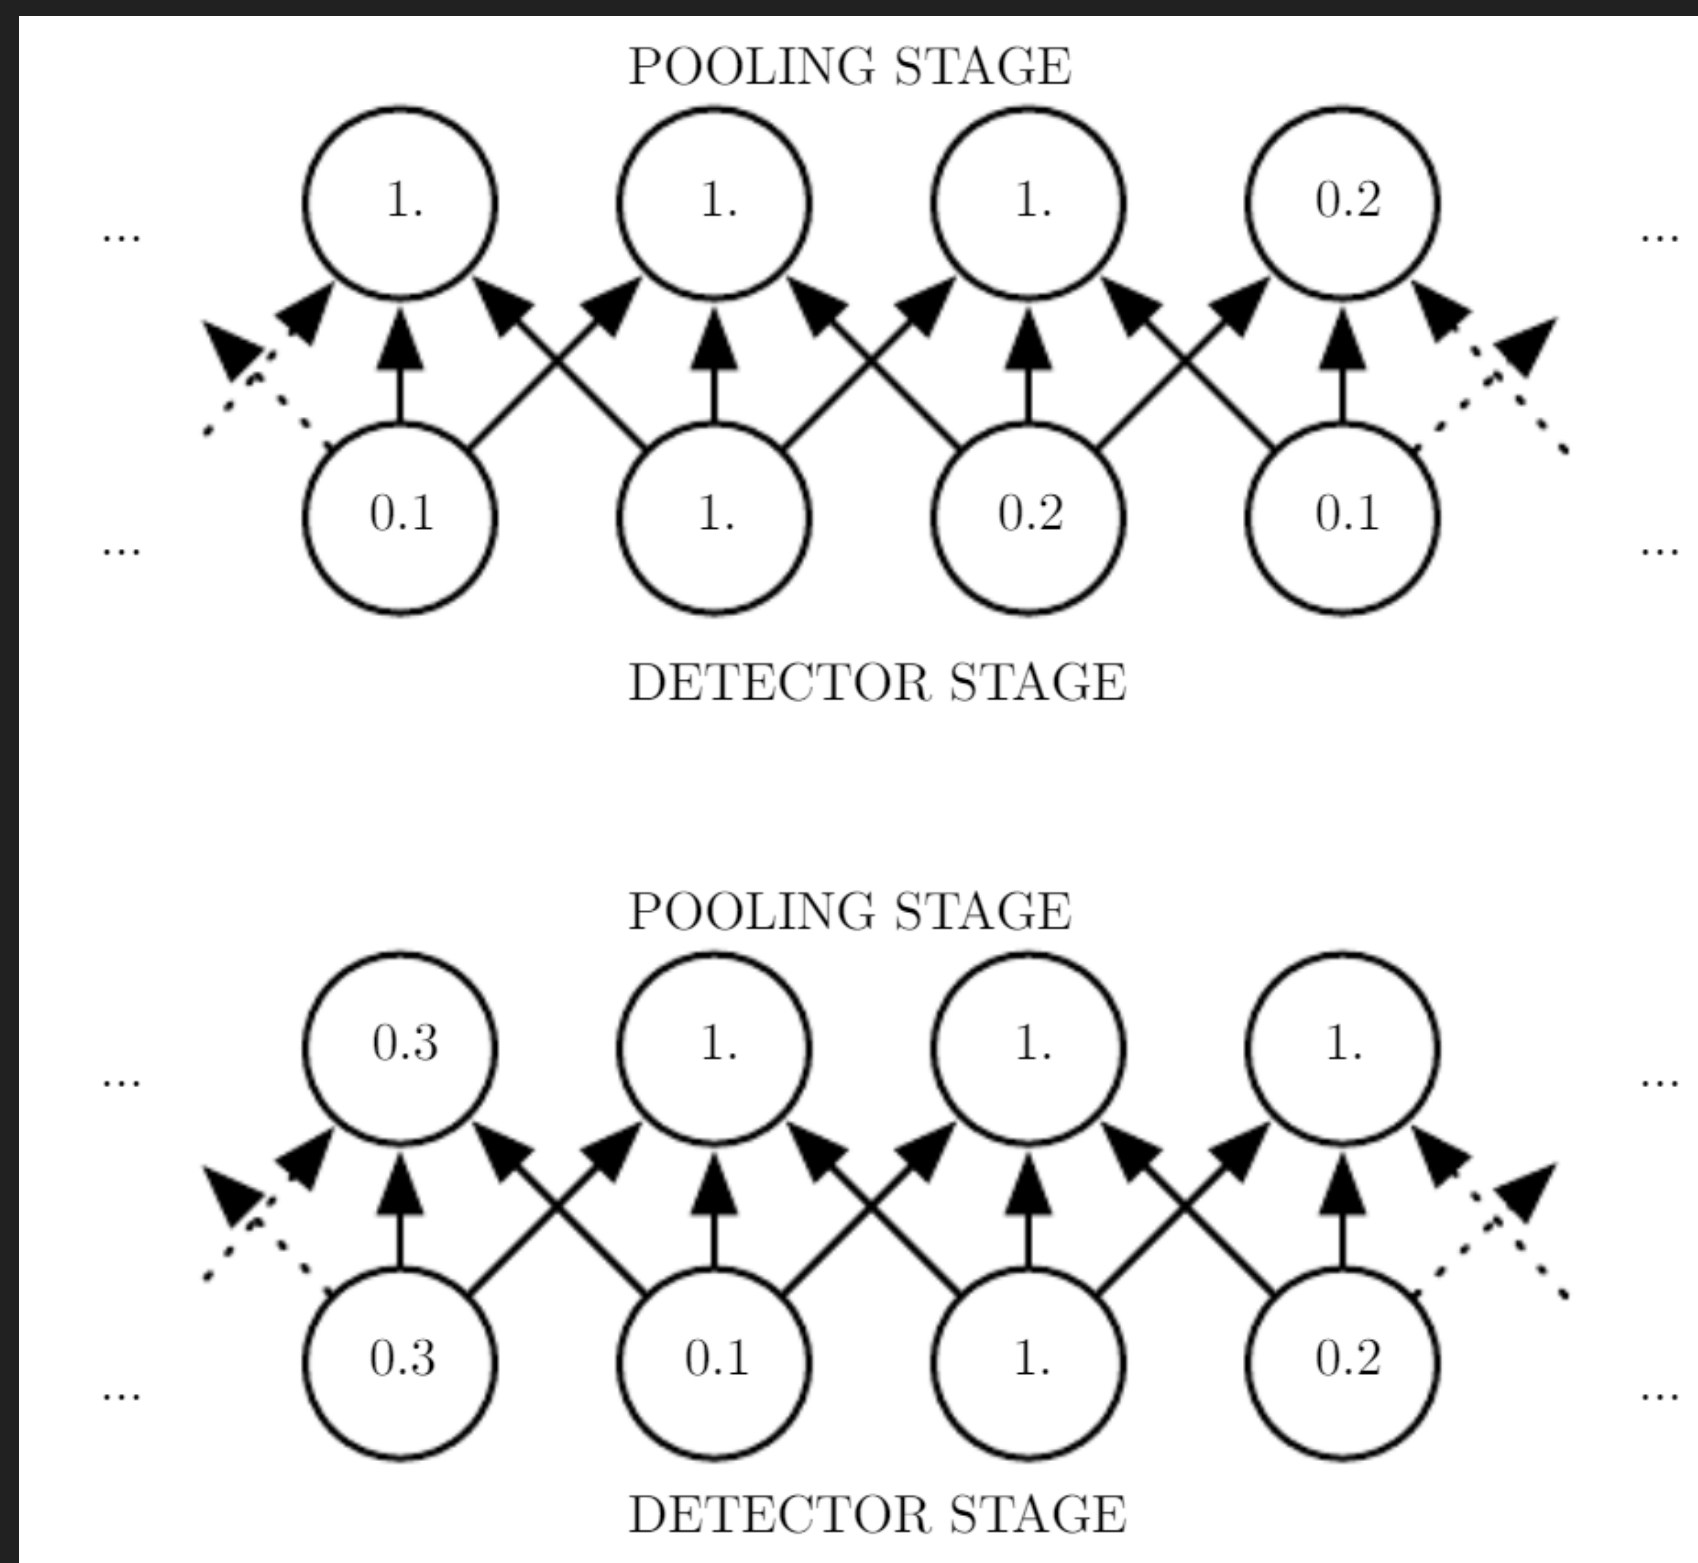


4	4
3	2

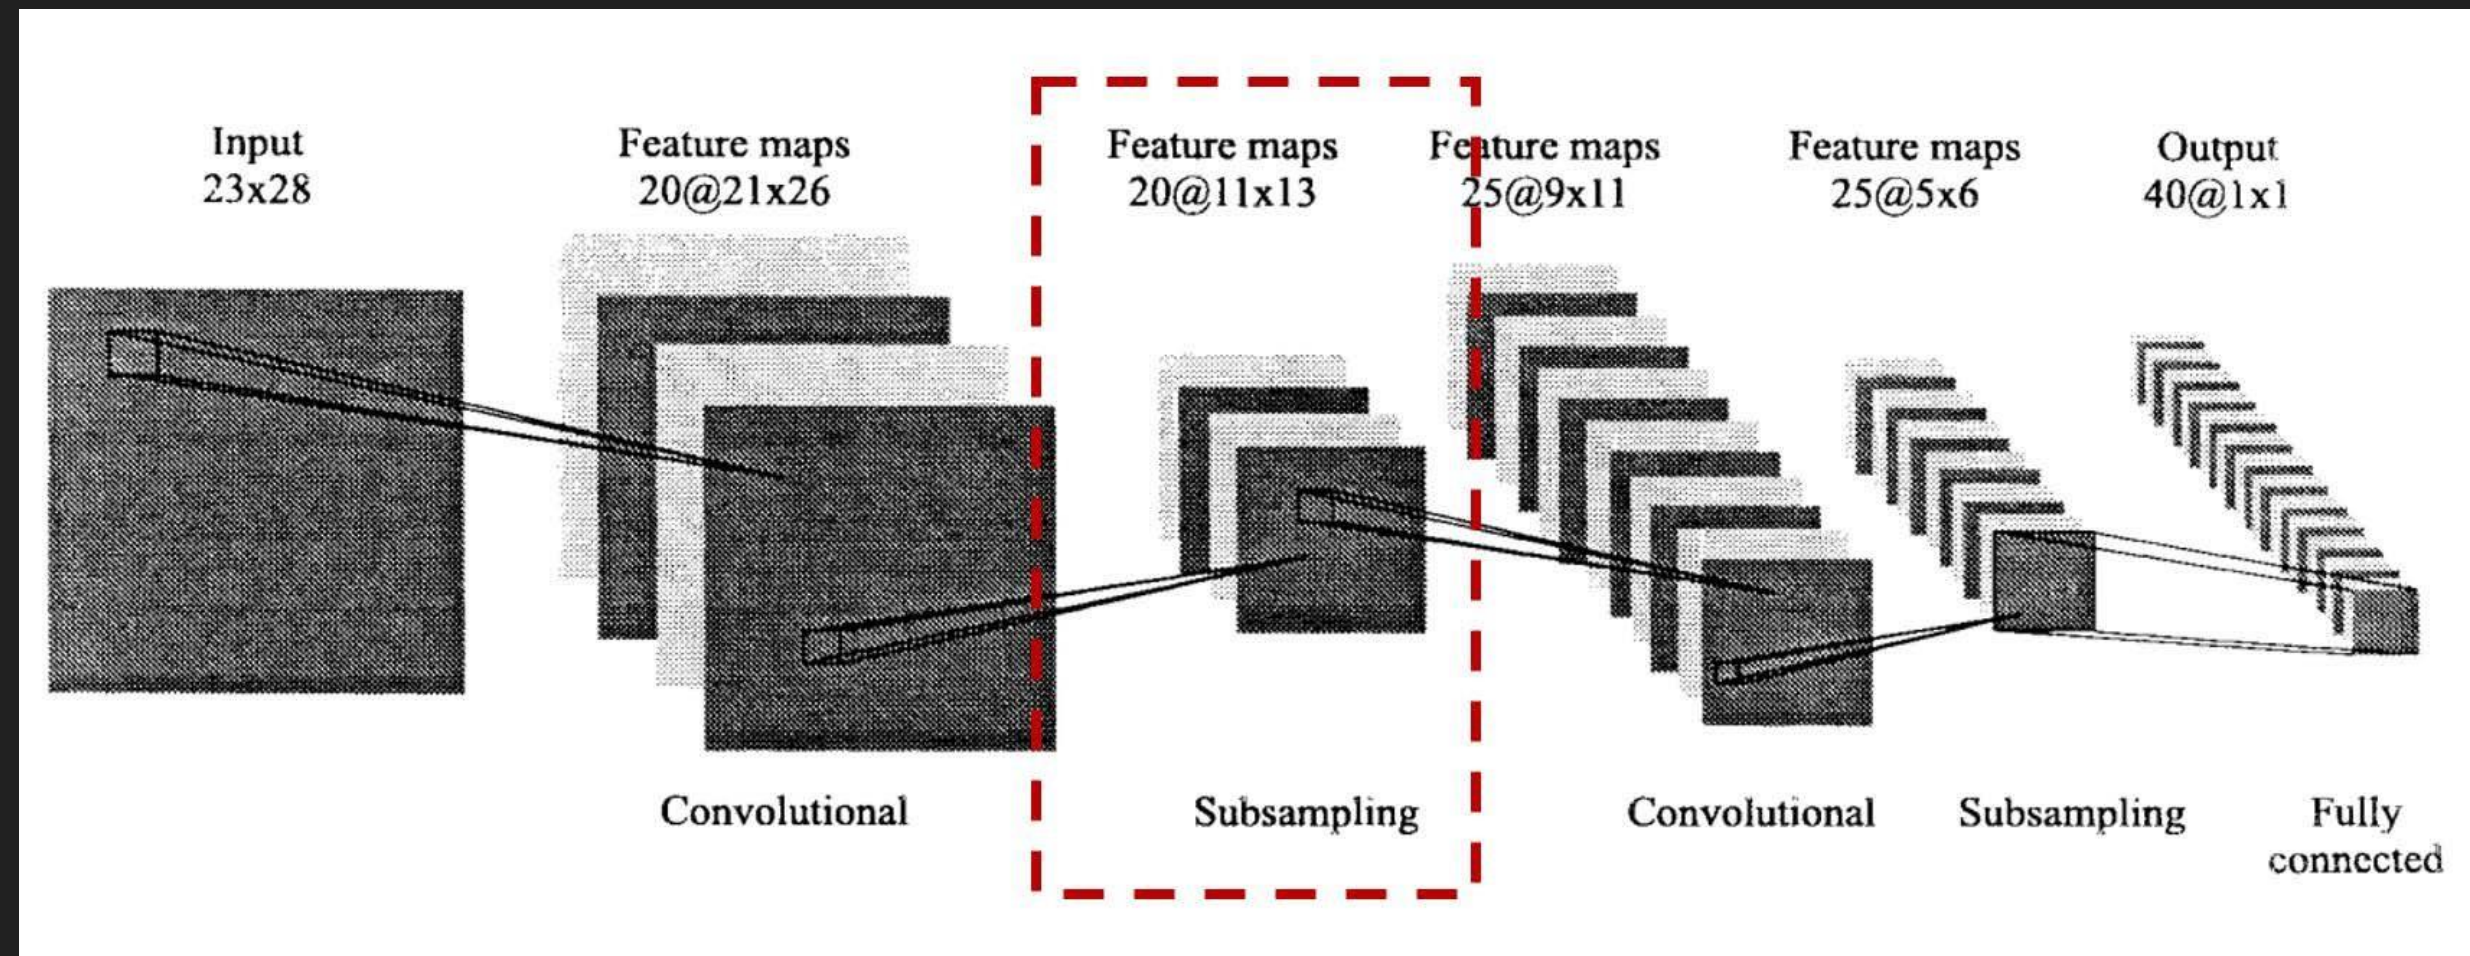
## POOLING WITH DOWNSAMPLING



## POOLING WITH NO DOWNSAMPLING



# LAYER






## ARCHITECTURE

- No one-size-fits-all approach to architectures
- Carefully consider input data including dimensions, size and channels
- Consider other parameters, such as pooling and kernel receptive field sizes
- Don't necessarily have to re-invent the wheel - other existing model architectures might be useful
- Can use other evolutionary algorithms to help determine right architectures and hyperparameters




# EXAMPLE ARCHITECTURES

## LeNet




# EXAMPLE ARCHITECTURES

## AlexNet




# EXAMPLE ARCHITECTURES

## VGG




# EXAMPLE ARCHITECTURES

## GoogLeNet




# EXAMPLE ARCHITECTURES

## ResNet



## RESIDUAL CONNECTIONS


- Add input of block to output.
- Creates clear path of unmodified gradients through all layers.







# PADDING

- Adds dummy data to the input
- As filter receptive fields are applied, the spatial dimensions can decrease faster than desired
- Padding allows spatial dimensions to be maintained and for filters to be applied in situations where it would exceed the input dimensions



## PADDING




## REGULARISATION

- A central goal of machine learning is making an algorithm perform well on new, unseen data points
- In ANNs and deep learning, several strategies exist for regularisation, including
  - Parameter sharing
  - Dataset augmentation
  - Dropout


## DROPOUT

- During each training pass, randomly remove a fraction of neural connections
- Adds noise to hidden units, particularly in the DNN classifier
- Inspired by sexual reproduction: half genes from both parents plus a small amount of random mutation results in offspring

# DROPOUT



(a) Standard Neural Net



(b) After applying dropout.

## BATCH NORMALISATION

- Batch normalisation reduces the the internal covariate shift
- Implicitly regularizes the model due to the noise in the batch estimates for mean and variance







## SUMMARY

- CNNs are inspired by the visual processing pathways of the brain
- Perform convolutions between input and weight tensors
- Pooling can downsample and provide translation invariance
- A DNN classifier maps CNN-produced representational features to outputs

